

Assignment05 - Amazon Fine Food Reviews Analysis_Logistic Regression

June 5, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [2]: # using SQLite Table to read data.
db_path = '/home/monodeepdas112/Datasets/amazon-fine-food-reviews/database.sqlite'
# db_path = '/home/monodeepdas112/Datasets/AmazonFineFoodReviews/database.sqlite'
con = sqlite3.connect(db_path)

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
```

```

# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    else:
        return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (500000, 10)

```

Out[2]:
   Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

   HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

   Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [3]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [4]: print(display.shape)
display.head()

```

(80668, 7)

```
Out [4]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ETO	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ETO	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBE1U	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [5]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [5]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	

	Score	Text	COUNT(*)
80638	5	I was recommended to try green tea extract to ...	5

```
In [6]: display['COUNT(*)'].sum()
```

```
Out [6]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [7]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [7]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	

3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [8]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)
```

```
In [9]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape
```

```
Out[9]: (348262, 10)
```

```
In [10]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[10]: 69.6524

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [11]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out[11]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0		3	1	5	1224892800
1		3	2	4	1212883200

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [12]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [13]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
```

```
final['Score'].value_counts()
```

```
(348260, 10)
```

```
Out[13]: 1    293516
0     54744
Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [14]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and I

=====

I've purchased both the Espresso Espresso (classic) and the 100% Arabica. My vote is defin

=====

This is a great product. It is very healthy for all of our dogs, and it is the first food that

=====

I find everything I need at Amazon so I always look there first. Chocolate tennis balls for a t

=====

```
In [15]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
```

```

sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)

```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and

```

In [16]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)

```

This book was purchased as a birthday gift for a 4 year old boy. He squealed with delight and

I've purchased both the Espressione Espresso (classic) and the 100% Arabica. My vote is defin

This is a great product. It is very healthy for all of our dogs, and it is the first food that

I find everything I need at Amazon so I always look there first. Chocolate tennis balls for a

```

In [17]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)

```



```

phrase = re.sub(r"\ 're", " are", phrase)
phrase = re.sub(r"\ 's", " is", phrase)
phrase = re.sub(r"\ 'd", " would", phrase)
phrase = re.sub(r"\ 'll", " will", phrase)
phrase = re.sub(r"\ 't", " not", phrase)
phrase = re.sub(r"\ 've", " have", phrase)
phrase = re.sub(r"\ 'm", " am", phrase)
return phrase

```

```

In [18]: sent_1500 = decontracted(sent_1500)
print(sent_1500)
print("="*50)

```

This is a great product. It is very healthy for all of our dogs, and it is the first food that
=====

```

In [19]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
print(sent_0)

```

This book was purchased as a birthday gift for a year old boy. He squealed with delight and h

```

In [20]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
print(sent_1500)

```

This is a great product It is very healthy for all of our dogs and it is the first food that t

```

In [21]: # https://gist.github.com/sebleier/554280
# we are removing the words from the stop words list: 'no', 'nor', 'not'
# <br /><br /> ==> after the above steps, we are getting "br br"
# we are including them into stop words list
# instead of <br /> if we have <br/> these tags would have reuvmoved in the 1st step

stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourself',
"you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 't
'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'h
'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as
'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through
'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'o
'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'a
'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too
's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n
've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't"

```

```
"hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
"mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
'won', "won't", 'wouldn', "wouldn't"])
```

```
In [22]: # Combining all the above students
from tqdm import tqdm
preprocessed_reviews = []
# tqdm is for printing the status bar
for sentence in tqdm(final['Text'].values):
    sentence = re.sub(r"http\S+", "", sentence)
    sentence = BeautifulSoup(sentence, 'lxml').get_text()
    sentence = decontracted(sentence)
    sentence = re.sub("\S*\d\S*", "", sentence).strip()
    sentence = re.sub('[^A-Za-z]+', ' ', sentence)
    # https://gist.github.com/sebleier/554280
    sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
    preprocessed_reviews.append(sentence.strip())
```

100%|| 348260/348260 [02:42<00:00, 2137.14it/s]

```
In [23]: preprocessed_reviews[1500]
```

```
Out[23]: 'great product healthy dogs first food love eat helped older dog lose weight year old'
```

[3.2] Preprocessing Review Summary

```
In [24]: # ## Similarly you can do preprocessing for review summary also.
# # Combining all the above students
# from tqdm import tqdm
# preprocessed_summary = []
# # tqdm is for printing the status bar
# for sentence in tqdm(final['Summary'].values):
#     sentence = re.sub(r"http\S+", "", sentence)
#     sentence = BeautifulSoup(sentence, 'lxml').get_text()
#     sentence = decontracted(sentence)
#     sentence = re.sub("\S*\d\S*", "", sentence).strip()
#     sentence = re.sub('[^A-Za-z]+', ' ', sentence)
#     # https://gist.github.com/sebleier/554280
#     sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
#     preprocessed_summary.append(sentence.strip())
```

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```
In [25]: # #BoW
# count_vect = CountVectorizer() #in scikit-learn
# count_vect.fit(preprocessed_reviews)
```

```
# print("some feature names ", count_vect.get_feature_names()[:10])
# print('='*50)

# final_counts = count_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_counts))
# print("the shape of out text BOW vectorizer ",final_counts.get_shape())
# print("the number of unique words ", final_counts.get_shape()[1])
```

5.2 [4.2] Bi-Grams and n-Grams.

In [26]: # #bi-gram, tri-gram and n-gram

```
# #removing stop words like "not" should be avoided before building n-grams
# # count_vect = CountVectorizer(ngram_range=(1,2))
# # please do read the CountVectorizer documentation http://scikit-learn.org/stable/m

# # you can choose these numebrs min_df=10, max_features=5000, of your choice
# count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
# final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_bigram_counts))
# print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_bigr
```

5.3 [4.3] TF-IDF

In [27]: # tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)

```
# tf_idf_vect.fit(preprocessed_reviews)
# print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_na
# print('='*50)

# final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_tf_idf))
# print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_tf_i
```

5.4 [4.4] Word2Vec

In [28]: # # Using Google News Word2Vectors

```
# # in this project we are using a pretrained model by google
# # its 3.3G file, once you load this into your memory
# # it occupies ~9Gb, so please do this step only if you have >12G of ram
# # we will provide a pickle file wich contains a dict ,
# # and it contains all our courpus words as keys and model[word] as values
# # To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# # from https://drive.google.com/file/d/0BYXkCwpI5KDYNlNUTTlSS21pQmM/edit
# # it's 1.9GB in size.
```

```

# # http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# # you can comment this whole cell
# # or change these variable according to your need

# is_your_ram_gt_16g=True
# want_to_use_google_w2v = False
# want_to_train_w2v = True

# path_to_word2vec = '/home/monodeepdas112/Datasets/GoogleNews-vectors-negative300.bi

# if want_to_train_w2v:

#     # Train your own Word2Vec model using your own text corpus
#     i=0
#     list_of_sentences=[]
#     for sentence in preprocessed_reviews:
#         list_of_sentences.append(sentence.split())

#     # min_count = 5 considers only words that occurred atleast 5 times
#     w2v_model=Word2Vec(list_of_sentences,min_count=5,size=100, workers=4)
#     print(w2v_model.wv.most_similar('great'))
#     print('='*50)
#     print(w2v_model.wv.most_similar('worst'))

# elif want_to_use_google_w2v and is_your_ram_gt_16g:
#     if os.path.isfile(path_to_word2vec):
#         print('Preparing to load pre-trained Word2Vec model !')
#         w2v_model=KeyedVectors.load_word2vec_format(path_to_word2vec, binary=True)
#         print('Successfully loaded model into memory !!')
#         print('Words similar to "similar" : ', w2v_model.wv.most_similar('great'))
#         print('Words similar to "worst" : ',w2v_model.wv.most_similar('worst'))
#     else:
#         print("you don't have google's word2vec file, keep want_to_train_w2v = True

In [29]: # w2v_words = list(w2v_model.wv.vocab)
# print("number of words that occurred minimum 5 times ",len(w2v_words))
# print("sample words ", w2v_words[0:50])

```

5.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```

In [30]: # # average Word2Vec
# # compute average word2vec for each review.
# sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
# for sent in tqdm(list_of_sentences): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
#     cnt_words = 0; # num of words with a valid vector in the sentence/review

```

```

#         for word in sent: # for each word in a review/sentence
#             if word in w2v_words:
#                 vec = w2v_model.wv[word]
#                 sent_vec += vec
#                 cnt_words += 1
#             if cnt_words != 0:
#                 sent_vec /= cnt_words
#             sent_vectors.append(sent_vec)
# print(len(sent_vectors))
# print(len(sent_vectors[0]))

```

[4.4.1.2] TFIDF weighted W2v

```

In [31]: # # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
# model = TfidfVectorizer()
# tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# # we are converting a dictionary with word as a key, and the idf as a value
# dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

```

```

In [32]: # # TF-IDF weighted Word2Vec
# tfidf_feat = model.get_feature_names() # tfidf words/col-names
# # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfi

# tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this
# row=0;
# for sent in tqdm(list_of_sentence): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length
#     weight_sum = 0; # num of words with a valid vector in the sentence/review
#     for word in sent: # for each word in a review/sentence
#         if word in w2v_words and word in tfidf_feat:
#             vec = w2v_model.wv[word]
#             #
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
#             # to reduce the computation we are
#             # dictionary[word] = idf value of word in whole corpus
#             # sent.count(word) = tf value of word in this review
#             tf_idf = dictionary[word]*(sent.count(word)/len(sent))
#             sent_vec += (vec * tf_idf)
#             weight_sum += tf_idf
#     if weight_sum != 0:
#         sent_vec /= weight_sum
#     tfidf_sent_vectors.append(sent_vec)
#     row += 1

```

6 [5] Assignment 5: Apply Logistic Regression

Apply Logistic Regression on these feature sets

SET 1:Review text, preprocessed one converted into vectors

- SET 2: Review text, preprocessed one converted into vectors
 - SET 3: Review text, preprocessed one converted into vectors
 - SET 4: Review text, preprocessed one converted into vectors

- Find the best hyper parameter which will give the maximum <https://www.appliedaicom.com>
- Find the best hyper parameter using k-fold cross validation or simple cross validation data
- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task

- Get the weights W after fit your model with the data X i.e Train data.
- Add a noise to the X ($X' = X + e$) and get the new data set X' (if X is a sparse matrix, $X.data += e$)
- Fit the model again on data X' and get the weights W'
- Add a small eps value(to eliminate the divisible by zero error) to W and W' i.e $W = W + 10^{-6}$ and $W' = W' + 10^{-6}$
- Now find the % change between W and W' ($| (W - W') / (W) | * 100$)
- Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden rise in it
- Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there is sudden rise
- Print the feature names whose % change is more than a threshold x (in our example 10)

- Calculate sparsity on weight vector obtained after using L1 regularization

- Get top 10 important features for both positive and negative classes separately.

- Feature engineering

```

<li>To increase the performance of your model, you can also experiment with with feature engine
    <ul>
        <li>Taking length of reviews as another feature.</li>
        <li>Considering some features from review summary as well.</li>
    </ul>
</li>
<br>
<li><strong>Representation of results</strong>
    <ul>
        <li>You need to plot the performance of model both on train data and cross validation data for
        <img src='train_cv_auc.JPG' width=300px></li>
        <li>Once after you found the best hyper parameter, you need to train your model with it, and f
        <img src='train_test_auc.JPG' width=300px></li>
        <li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
        <img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
        <li>You need to summarize the results at the end of the notebook, summarize it in the table for
        <img src='summary.JPG' width=400px>
    </li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

7 Applying Logistic Regression

```

In [33]: #Getting the necessary imports and function definations
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score
from scipy.stats import uniform
from sklearn.model_selection import train_test_split
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import StratifiedKFold
import pprint
from sklearn.pipeline import Pipeline

```

```

import os.path
import pickle
import pprint
from sklearn import preprocessing
import math

import warnings
warnings.filterwarnings('ignore')

```

7.0.1 [5.0.0] Splitting up the Dataset into D_train and D_test

```
In [34]: num_data_points = 50000
```

```
In [35]: Dx_train, Dx_test, Dy_train, Dy_test = train_test_split(preprocessed_reviews[:num_data_points],
```

```
In [36]: prettytable_data = []
```

7.0.2 [5.0.1] Defining some functions to increase code reusability and readability

```

In [37]: '''Creating Custom Vectorizers for TFIDF - W2Vec and Avg - W2Vec'''
class Tfidf_W2Vec_Vectorizer(object):
    def __init__(self, w2vec_model):
        if(w2v_model is None):
            raise Exception('Word 2 Vector model passed to Tfidf_W2Vec Vectorizer is None')
        self.tfidf = TfidfVectorizer(max_features=300)
        self.dictionary = None
        self.tfidf_feat = None

        self.word2vec = w2vec_model

    def fit(self, X):
        '''X : list'''
        #Initializing the TFIDF Vectorizer
        self.tfidf.fit_transform(X)
        # we are converting a dictionary with word as a key, and the idf as a value
        self.dictionary = dict(zip(self.tfidf.get_feature_names(), list(self.tfidf.idf)))
        self.tfidf_feat = self.tfidf.get_feature_names()

        return self

    def transform(self, X):
        '''X : list'''
        return np.array([
            np.mean([self.word2vec[w] * self.dictionary[word]*(X.count(word)/len(X))
                    for w in words if w in self.word2vec and w in self.tfidf_feat],
                    axis=0)
            for words in X
        ])

```



```

class Avg_W2Vec_Vectorizer(object):
    def __init__(self, w2vec_model):
        if(w2v_model is None):
            raise Exception('Word 2 Vector model passed to Avg_W2Vec Vectorizer is None')
        self.word2vec = w2vec_model

    def fit(self, X):
        return self

    def transform(self, X):
        '''X : list'''
        return np.array([
            np.mean([self.word2vec[w] for w in words if w in self.word2vec]
                    or [np.zeros(300)], axis=0)
            for words in X
        ])

In [38]: def get_vectorizer(vectorizer, train, W2V_model=None):
    if(vectorizer=='BOW'):
        vectorizer = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=1000)
    if(vectorizer=='TFIDF'):
        vectorizer = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=1000)
    if(vectorizer=='TFIDF-W2Vec'):
        vectorizer = Tfidf_W2Vec_Vectorizer(W2V_model)
    if(vectorizer=='Avg-W2Vec'):
        vectorizer = Avg_W2Vec_Vectorizer(W2V_model)

    vectorizer.fit(train)
    return vectorizer

In [39]: '''Perform Simple Cross Validation'''
def perform_hyperparameter_tuning(X, Y, vectorizer, vec_name, penalty, results_path,
    #If the pandas dataframe with the hyperparameter info exists then return it

    if(retrain==False):
        # If Cross Validation results exists then return them
        if(os.path.exists(results_path)):
            return pd.read_csv(results_path)
        else:
            # If no data exists but retrain=False then mention accordingly
            print('Retrain is set to be False but no Cross Validation Results DataFrame')
    else:
        # else perform hyperparameter tuning
        print('Performing Hyperparameter Tuning...\n')
        # regularization parameter
        c = [0.0001, 0.001, 0.01, 0.05, 0.1, 0.5, 1, 5, 10, 50, 100, 500, 1000, 5000,
        c.sort()
        hyperparameters = {

```

```

        'logistic__penalty' : penalty,
        'logistic__C' : c
    }

    penalties = []
    C_values = []

    train_scores = []
    test_scores = []

    train_mean_score = []
    test_mean_score = []

    # Initializing KFold
    skf = StratifiedKFold(n_splits=3)
    X = np.array(X)
    Y = np.array(Y)

    saver = 0 # This is a counter variable that saves
    for penalty in hyperparameters['logistic__penalty']:
        for reg_param in hyperparameters['logistic__C']:

            #Performing Cross Validation
            for train_index, test_index in skf.split(X, Y):
                Dx_train, Dx_cv = X[train_index], X[test_index]
                Dy_train, Dy_cv = Y[train_index], Y[test_index]

                #Initializing the Vectorizer
                vectorizer = get_vectorizer(vectorizer, Dx_train.tolist(), W2V_model)

                #Transforming the data to features
                x_train = vectorizer.transform(Dx_train.tolist())
                x_cv = vectorizer.transform(Dx_cv.tolist())

                #Initializing the LR model
                log_reg_model = LogisticRegression(penalty=penalty, C=reg_param, max_iter=1000)

                # Fit the model
                log_reg_model.fit(x_train, Dy_train)

                #Prediction
                train_results = log_reg_model.predict_proba(x_train)
                cv_results = log_reg_model.predict_proba(x_cv)

            try:
                train_score = roc_auc_score(Dy_train, train_results[:, 1])
                test_score = roc_auc_score(Dy_cv, cv_results[:, 1])

```

```

        #storing the results to form a dataframe
        train_scores.append(train_score)
        test_scores.append(test_score)

    except Exception as e:
        print('Error Case : ', e)
        print(('Actual, Predicted'))
        [print((Dy_cv[i], cv_results[i, 1])) for i in range(len(Dy_cv))]

    print('CV iteration : C={0}, solver={1}, train_score={2}, test_score={3}'
          .format(reg_param, 'saga', train_score, test_score))

    train_mean_score.append(sum(train_scores)/len(train_scores))
    test_mean_score.append(sum(test_scores)/len(test_scores))

    penalties.append(penalty)
    C_values.append(reg_param)

    print('C={0}, penalty={1}, solver="saga", train_score={2}, test_score={3}'
          .format(reg_param, penalty, sum(train_scores)/len(train_scores)))

    saver += 1

    if(saver==5):
        # after every period of 100 iterations keep saving the parameters
        # so as to avoid data loss in case of system crash
        with open('saved_temp_data/{0}_penalties.pkl'.format(vec_name), 'w') as file:
            pickle.dump(penalties, file)
        with open('saved_temp_data/{0}_reg_params.pkl'.format(vec_name), 'w') as file:
            pickle.dump(C_values, file)
        with open('saved_temp_data/{0}_train_mean_score.pkl'.format(vec_name), 'w') as file:
            pickle.dump(train_mean_score, file)
        with open('saved_temp_data/{0}_test_mean_score.pkl'.format(vec_name), 'w') as file:
            pickle.dump(test_mean_score, file)
        saver = 0

    train_scores = []
    test_scores = []

try:
    # Attempting to load saved data
    # Load data from the saved files
    with open('saved_temp_data/{0}_penalties.pkl'.format(vec_name), 'rb') as file:
        penalties = pickle.load(file)
    with open('saved_temp_data/{0}_reg_params.pkl'.format(vec_name), 'rb') as file:
        C_values = pickle.load(file)
    with open('saved_temp_data/{0}_train_mean_score.pkl'.format(vec_name), 'rb') as file:
        train_mean_score = pickle.load(file)
    with open('saved_temp_data/{0}_test_mean_score.pkl'.format(vec_name), 'rb') as file:
        test_mean_score = pickle.load(file)

```

```

        test_mean_score = pickle.load(file)
    except Exception as ex:
        print('Failed to load saved data from temp files')
    # Creating a DataFrame from the saved data for visualization
    results_df = pd.DataFrame({'C' : C_values, 'penalty' : penalties,
                              'solver' : ['saga' for i in C_values], 'train_score' : train_scores,
                              'test_score' : test_mean_score})

    try:
        # Attempting to remove the temporary files
        os.remove('saved_temp_data/{0}_penalties.pkl'.format(vec_name))
        os.remove('saved_temp_data/{0}_reg_params.pkl'.format(vec_name))
        os.remove('saved_temp_data/{0}_train_mean_score.pkl'.format(vec_name))
        os.remove('saved_temp_data/{0}_test_mean_score.pkl'.format(vec_name))
    except Exception as e:
        print('Error occurred while attempting to remove the temporary files')

    #writing the results to csv after performing hyperparameter tuning
    try:
        results_df.to_csv(results_path)
    except Exception as ex:
        print(str(ex), "\nError occured while converting DataFrame to CSV after c")
    return results_df

```

```

In [40]: def analyse_results(df):
    # plotting error curves
    fig = plt.figure()
    ax = fig.gca()

    plt.plot([math.log10(i) for i in df.C.tolist()], df.test_score.tolist(), '-o', c='b')
    plt.plot([math.log10(i) for i in df.C.tolist()], df.train_score.tolist(), '-o', c='r')
    plt.grid(True)
    plt.xlabel('log10 of Hyperparameter C = 1/alpha')
    plt.ylabel('Area Under ROC Curve')
    plt.title('AUC ROC Curve for Logistic Regression')
    plt.legend(loc='best')
    plt.show()

    # return the best parameters
    mmax = 0
    ind_max = 0
    for index, row in df.iterrows():
        if(row['test_score']>mmax):
            mmax=row['test_score']
            ind_max = index

    best_params = {

```

```

        'logistic__C':df.loc[ind_max, 'C'],
        'logistic__penalty':df.loc[ind_max, 'penalty'],
        'logistic__solver':df.loc[ind_max, 'solver']
    }

    return best_params

In [41]: def retrain_with_best_params(data, labels, best_params, vec_name, model_path, word2vec):
    if(os.path.exists(model_path)):
        print('Loading Model....')
        with open(model_path, 'rb') as input_file:
            clf = pickle.load(input_file)
    else:
        clf = LogisticRegression(penalty=best_params['logistic__penalty'],
                                C = best_params['logistic__C'],
                                solver=best_params['logistic__solver'], max_iter=1000)
        print('Initializing Vectorizer')
        vectorizer = get_vectorizer(vectorizer=vec_name, train=data, W2V_model=word2vec)
        print('Training Model....')
        clf.fit(vectorizer.transform(data), np.array(labels))

        print('Saving Trained Model....')
        with open(model_path,'wb') as file:
            pickle.dump(clf,file)
    return clf

In [42]: def plot_confusion_matrix(model, data, labels, dataset_label):
    pred = model.predict(data)
    conf_mat = confusion_matrix(labels, pred)

    strings = strings = np.asarray([['TN = ', 'FP = '],
                                    ['FN = ', 'TP = ']])

    labels = (np.asarray(["{0}{1}".format(string, value)
                           for string, value in zip(strings.flatten(),
                                                    conf_mat.flatten())])
              ).reshape(2, 2)

    fig, ax = plt.subplots()
    ax.set(xlabel='Predicted', ylabel='Actual', title='Confusion Matrix : {0}'.format(dataset_label))
    sns.heatmap(conf_mat, annot=labels, fmt="", cmap='YlGnBu', ax=ax)
    ax.set_xlabel('Predicted')
    ax.set_ylabel('Actual')
    ax.set_xticklabels(['False', 'True'])
    ax.set_yticklabels(['False', 'True'])
    plt.show()

In [43]: def plot_AUC_ROC(model, vectorizer, Dx_train, Dx_test, Dy_train, Dy_test):

```

```

#predicting probability of Dx_test, Dx_train
test_score = model.predict_proba(vectorizer.transform(Dx_test))
train_score = model.predict_proba(vectorizer.transform(Dx_train))

#Finding out the ROC_AUC_SCORE
train_roc_auc_score = roc_auc_score(np.array(Dy_train), train_score[:, 1])
print('Area Under the Curve for Train : ', train_roc_auc_score)
test_roc_auc_score = roc_auc_score(np.array(Dy_test), test_score[:, 1])
print('Area Under the Curve for Test : ', test_roc_auc_score)

#Plotting with matplotlib.pyplot
#ROC Curve for D-train
train_fpr, train_tpr, thresholds = roc_curve(np.array(Dy_train), train_score[:, 1])
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))

# ROC Curve for D-test
test_fpr, test_tpr, thresholds = roc_curve(np.array(Dy_test), test_score[:, 1])
plt.plot(test_fpr, test_tpr, label="train AUC =" + str(auc(test_fpr, test_tpr)))

plt.legend()
plt.xlabel("FPR : False Positive Ratio")
plt.ylabel("TPF : True Positive Ratio")
plt.title("Area Under ROC Curve")
plt.show()

plot_confusion_matrix(model, vectorizer.transform(Dx_train), np.array(Dy_train),
plot_confusion_matrix(model, vectorizer.transform(Dx_test), np.array(Dy_test), 'T
return train_roc_auc_score, test_roc_auc_score

```

7.1 [5.1.0] Logistic Regression on BOW, SET 1

7.1.1 [5.1.0] Applying Logistic Regression with L1 regularization on BOW, SET 1

```

In [44]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/BOW_log_reg_results_l1.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='BOW', v
penalty=['l1'], results_path=csv_path, ret

# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)
# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l1.pkl'.format('BOW')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'BOW', model_

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = None, train=Dx_train, vectorizer='BOW')

# plotting AUC ROC

```

```

train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy)

# appending the data results
# prettytable_data.append(['BOW', 'LogisticRegression', 'L1', best_parameters['logist

```


Performing Hyperparameter Tuning...

```

CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
C=0.0001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
C=0.001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
CV iteration : C=0.01, solver=saga, train_score=0.8244418426264546, test_score=0.82065213504638
CV iteration : C=0.01, solver=saga, train_score=0.8247320630638473, test_score=0.8115282727301
CV iteration : C=0.01, solver=saga, train_score=0.8239450915217509, test_score=0.8342770795735
C=0.01, penalty=l1, solver="saga", train_score=0.8243729990706843, test_score=0.82215249578335
CV iteration : C=0.05, solver=saga, train_score=0.9019193121821065, test_score=0.8900017362818
CV iteration : C=0.05, solver=saga, train_score=0.9036854023941612, test_score=0.8929765991173
CV iteration : C=0.05, solver=saga, train_score=0.9015880912165825, test_score=0.8944236887819
C=0.05, penalty=l1, solver="saga", train_score=0.9023976019309501, test_score=0.89246734139374
CV iteration : C=0.1, solver=saga, train_score=0.9216839057870274, test_score=0.90467340608209
CV iteration : C=0.1, solver=saga, train_score=0.9250289150594848, test_score=0.90670204805057
CV iteration : C=0.1, solver=saga, train_score=0.9226615530566592, test_score=0.90743292424407
C=0.1, penalty=l1, solver="saga", train_score=0.9231247913010572, test_score=0.906269459458915
CV iteration : C=0.5, solver=saga, train_score=0.9386440024187235, test_score=0.91328142503136
CV iteration : C=0.5, solver=saga, train_score=0.9453715094188645, test_score=0.91367030058779
CV iteration : C=0.5, solver=saga, train_score=0.9414061945871383, test_score=0.91207727078307
C=0.5, penalty=l1, solver="saga", train_score=0.9418072354749087, test_score=0.913009665467411
CV iteration : C=1, solver=saga, train_score=0.9399987308443226, test_score=0.912962159020207
CV iteration : C=1, solver=saga, train_score=0.9469477978390353, test_score=0.9117054942359779
CV iteration : C=1, solver=saga, train_score=0.9427849911503553, test_score=0.9107327775422377
C=1, penalty=l1, solver="saga", train_score=0.9432438399445711, test_score=0.9118001435994741
CV iteration : C=5, solver=saga, train_score=0.9406982466599663, test_score=0.9121590846875297
CV iteration : C=5, solver=saga, train_score=0.9476213150152841, test_score=0.9093168463540602
CV iteration : C=5, solver=saga, train_score=0.9434832554898338, test_score=0.9093009338641114
C=5, penalty=l1, solver="saga", train_score=0.9439342723883614, test_score=0.9102589549685671
CV iteration : C=10, solver=saga, train_score=0.9407587294022335, test_score=0.912048412837887
CV iteration : C=10, solver=saga, train_score=0.9476632440498435, test_score=0.908975100123439
CV iteration : C=10, solver=saga, train_score=0.9435407456468976, test_score=0.909113628731798
C=10, penalty=l1, solver="saga", train_score=0.9439875730329915, test_score=0.9100457138977083
CV iteration : C=50, solver=saga, train_score=0.9407979911581177, test_score=0.911949581938260
CV iteration : C=50, solver=saga, train_score=0.9476800684230258, test_score=0.908692500024326
CV iteration : C=50, solver=saga, train_score=0.9435776899971288, test_score=0.908945611045034
C=50, penalty=l1, solver="saga", train_score=0.9440185831927574, test_score=0.9098625643358739
CV iteration : C=100, solver=saga, train_score=0.9408030619186761, test_score=0.91193750651398

```

CV iteration : C=100, solver=saga, train_score=0.9476835270920771, test_score=0.90865369453467
 CV iteration : C=100, solver=saga, train_score=0.943581661404869, test_score=0.908924799363666
 C=100, penalty=l1, solver="saga", train_score=0.9440227501385409, test_score=0.909838666804109
 CV iteration : C=500, solver=saga, train_score=0.9408071507689529, test_score=0.91193023781199
 CV iteration : C=500, solver=saga, train_score=0.947682427938777, test_score=0.908625850716571
 CV iteration : C=500, solver=saga, train_score=0.9435838742556762, test_score=0.90890815001857
 C=500, penalty=l1, solver="saga", train_score=0.9440244843211354, test_score=0.909821412849047
 CV iteration : C=1000, solver=saga, train_score=0.9408063007570675, test_score=0.9119297102449
 CV iteration : C=1000, solver=saga, train_score=0.9476834831259451, test_score=0.9086211612314
 CV iteration : C=1000, solver=saga, train_score=0.943584929389836, test_score=0.90890574641593
 C=1000, penalty=l1, solver="saga", train_score=0.9440249044242828, test_score=0.90981887263075
 CV iteration : C=5000, solver=saga, train_score=0.9408073412888583, test_score=0.9119239656256
 CV iteration : C=5000, solver=saga, train_score=0.9476831167415117, test_score=0.9086197543858
 CV iteration : C=5000, solver=saga, train_score=0.9435839035649584, test_score=0.9089054532936
 C=5000, penalty=l1, solver="saga", train_score=0.9440247871984427, test_score=0.90981639110171
 CV iteration : C=10000, solver=saga, train_score=0.9408075611195184, test_score=0.911920448511
 CV iteration : C=10000, solver=saga, train_score=0.9476830141538704, test_score=0.908618288921
 CV iteration : C=10000, solver=saga, train_score=0.9435859991786366, test_score=0.908904046306
 C=10000, penalty=l1, solver="saga", train_score=0.9440255248173418, test_score=0.9098142612467




```
{'logistic__C': 0.5, 'logistic__penalty': 'l1', 'logistic__solver': 'saga'}
Initializing Vectorizer
Training Model...
```




Saving Trained Model...

Retraining Vectorizer with Dx_train

Area Under the Curve for Train : 0.9373688316842859

Area Under the Curve for Test : 0.9122061345473661





[5.1.0.1] Calculating sparsity on weight vector obtained using L1 regularization on BOW, SET 1

```
In [45]: # # Please write all the code with proper documentation
# # fetching the coeff and calculating sparsity
weights = log_reg.coef_
elements = (weights.shape[0]*weights.shape[1])
zeroes = elements - np.count_nonzero(weights)
sparsity = zeroes/elements
print('Sparsity of the weight vector by L1 regularization is : ', sparsity)
```

Sparsity of the weight vector by L1 regularization is : 0.137

7.1.2 [5.1.0.2] Feature Importance on BOW, SET 1

```
In [46]: # Please write all the code with proper documentation
feature_names = vectorizer_obj.get_feature_names()
weights = np.reshape(log_reg.coef_,(log_reg.coef_.shape[1], log_reg.coef_.shape[0]))

# making a list of feature names along with their feature weights
features_with_weights = [(feature_names[i],weights[i]) for i in range(len(feature_names))]
features_with_weights.sort(key=lambda x : abs(x[1][0]), reverse=True)
```

[5.1.0.2.1] Top 10 important features of positive class from SET 1

```
In [47]: positive_weights = [i for i in features_with_weights if i[1][0]>=0]
print('Top 10 features of positive class with the feature names : ')
for i in positive_weights[:10]:
    print(i[0], ' : ',i[1][0])
```

Top 10 features of positive class with the feature names :

```
excellent : 1.9753053108460117
delicious : 1.8589530909799676
pleased : 1.8324072393701578
loves : 1.6731939586153268
perfect : 1.5793338333615692
awesome : 1.5033011438383237
keeps : 1.4020567949287117
best : 1.4003388186598489
wonderful : 1.3736884238797988
amazing : 1.3533707686471173
```

[5.1.0.2.2] Top 10 important features of negative class from SET 1

```
In [48]: negative_weights = [i for i in features_with_weights if i[1][0]<0]
print('Top 10 features of negative class with the feature names : ')
for i in negative_weights[:10]:
    print(i[0], ' : ',i[1][0])
```

Top 10 features of negative class with the feature names :

```
terrible : -1.906277472024034
awful : -1.7446309138069729
not buy : -1.5747627365624441
unfortunately : -1.511689812131779
not good : -1.4807794459805517
disappointed : -1.4336116407383688
horrible : -1.3764205956238866
money : -1.225686396423233
bland : -1.0690434464749512
ended : -0.929604547728539
```

[5.1.0.2.3] Performing perturbation test (multicollinearity check) on BOW, SET 1

```
In [49]: # Get the weights W after fit your model with the data X i.e Train data.
# Add a noise to the X ( $X' = X + e$ ) and get the new data set X' (if X is a sparse mat
# Fit the model again on data X' and get the weights W'
# Add a small eps value(to eliminate the divisible by zero error) to W and W i.e  $W=W+$ 
# Now find the % change between W and W' ( $| (W-W') / (W) | * 100$ )
# Calculate the 0th, 10th, 20th, 30th, ...100th percentiles, and observe any sudden r
# Ex: consider your 99th percentile is 1.3 and your 100th percentiles are 34.6, there
# Print the feature names whose % change is more than a threshold x(in our example it
```

```
In [50]: def pertubation_test(model, features, labels, features_with_weights):

    # saving initial model weights
    W = model.coef_

    # change the data slightly
    print('Adding noise to features...')
    noise = np.random.normal(0,0.1,features.shape[0]*features.shape[1])
    noise = np.reshape(noise, features.shape)
    features = features + noise

    # retraining the model with the new features formed by adding errors
    print('Retraining model...')
    model.fit(features, labels)

    # getting the model weights after retraining
    print('Calculating model weights percentage change...')
    _W = model.coef_
    epsilon_val = np.random.normal(0, 0.000005, model.coef_.shape[0]*model.coef_.shape[1])
    epsilon_val = np.reshape(epsilon_val, _W.shape)
    _W = _W + epsilon_val
    W = W + epsilon_val

    percent_change = abs((W-_W)/W)*100
```

```

#      0 : feature name
#      1 : percent change
feature_percent_change = [(features_with_weights[i][0], percent_change[0][i])for i in range(0, 100, 10)]

# sorting the data according to the weight values
feature_percent_change.sort(key=lambda x: x[1])

# calculating percentile values and displaying
print('Percentile values of some intervals')
for i in range(0, 100, 10):
    index = int((i/100)*(percent_change.shape[1]+1))
    print('{0} th percentile : '.format(i),feature_percent_change[index])

# printing those feature names with percent change difference above threshold
print('Percent Change above threshold = 2.5')
for i in range(1, 100):
    indexi = int((i/100)*(percent_change.shape[1]+1))
    indexi_ = int(((i-1)/100)*(percent_change.shape[1]+1))
    difference = abs(feature_percent_change[indexi][1]-feature_percent_change[indexi_][1])
    if(difference > 2.5): # taking threshold of 2.5
        print('\n{0} th percentile : '.format(i-1),feature_percent_change[indexi_][1])
        print('{0} th percentile : '.format(i),feature_percent_change[indexi][1])

```

In [51]: pertubation_test(log_reg, vectorizer_obj.transform(Dx_train), Dy_train, features_with_weights)

Adding noise to features...

Retraining model...

Calculating model weights percentage change...

Percentile values of some intervals

```

0 th percentile : ('keeps', 0.0)
10 th percentile : ('spice', 7.426537586476137)
20 th percentile : ('not eat', 21.118749524096533)
30 th percentile : ('tea bags', 31.57523264918191)
40 th percentile : ('bite', 41.16511261235236)
50 th percentile : ('yes', 52.37539559627573)
60 th percentile : ('least', 66.50152315798424)
70 th percentile : ('apple', 91.66743791896737)
80 th percentile : ('house', 113.70426985541637)
90 th percentile : ('food', 299628.3570150673)

```

Percent Change above threshold = 2.5

```

63 th percentile : ('worked', 72.13567526873742)
64 th percentile : ('milk', 75.26485866456457)

```

```

65 th percentile : ('fun', 76.49652010742746)
66 th percentile : ('touch', 79.67994081218664)

```

```

66 th percentile : ('touch', 79.67994081218664)

```

67 th percentile : ('opened', 83.96846207049516)
 67 th percentile : ('opened', 83.96846207049516)
 68 th percentile : ('cheaper', 86.85848087514947)
 68 th percentile : ('cheaper', 86.85848087514947)
 69 th percentile : ('serving', 90.12445915407207)
 70 th percentile : ('apple', 91.66743791896737)
 71 th percentile : ('varieties', 95.16785731890529)
 71 th percentile : ('varieties', 95.16785731890529)
 72 th percentile : ('hooked', 98.91544871834057)
 79 th percentile : ('easy make', 102.28756771874497)
 80 th percentile : ('house', 113.70426985541637)
 80 th percentile : ('house', 113.70426985541637)
 81 th percentile : ('mean', 124.1882682730498)
 81 th percentile : ('mean', 124.1882682730498)
 82 th percentile : ('room', 142.97938447766813)
 82 th percentile : ('room', 142.97938447766813)
 83 th percentile : ('soft', 165.07956862597618)
 83 th percentile : ('soft', 165.07956862597618)
 84 th percentile : ('cocoa', 217.64991389179227)
 84 th percentile : ('cocoa', 217.64991389179227)
 85 th percentile : ('dogs love', 306.09064033396805)
 85 th percentile : ('dogs love', 306.09064033396805)
 86 th percentile : ('ok', 368.8947319406195)
 86 th percentile : ('ok', 368.8947319406195)
 87 th percentile : ('children', 425.15531719705876)
 87 th percentile : ('children', 425.15531719705876)
 88 th percentile : ('excellent', 985.4557058363861)
 88 th percentile : ('excellent', 985.4557058363861)
 89 th percentile : ('due', 35268.29847831288)
 89 th percentile : ('due', 35268.29847831288)
 90 th percentile : ('food', 299628.3570150673)
 90 th percentile : ('food', 299628.3570150673)

```

91 th percentile : ('best', 511176.085627835)

91 th percentile : ('best', 511176.085627835)
92 th percentile : ('daughter', 746075.4445921665)

92 th percentile : ('daughter', 746075.4445921665)
93 th percentile : ('always', 987063.9099715654)

93 th percentile : ('always', 987063.9099715654)
94 th percentile : ('though', 1521587.0195252018)

94 th percentile : ('though', 1521587.0195252018)
95 th percentile : ('crazy', 2276999.184031091)

95 th percentile : ('crazy', 2276999.184031091)
96 th percentile : ('flavors', 2903741.576993312)

96 th percentile : ('flavors', 2903741.576993312)
97 th percentile : ('thin', 4680354.412447514)

97 th percentile : ('thin', 4680354.412447514)
98 th percentile : ('bags', 7390505.314619637)

98 th percentile : ('bags', 7390505.314619637)
99 th percentile : ('eggs', 23824140.225007847)

```

These are some of the features showing multicollinearity and hence we cant give feature importance or interpretability.

7.1.3 [5.1.2] Applying Logistic Regression with L2 regularization on BOW, SET 1

```

In [52]: # Please write all the code with proper documentation
        csv_path = 'saved_models/Assignment5/BOW_log_reg_results_l2.csv'
        cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='BOW', v
                                                    penalty=['l2'], results_path=csv_path, ret

        # Analysing best parameters
        best_parameters = analyse_results(cv_results)

        # retraining the model with best parameters
        model_path = 'saved_models/Assignment5/{0}_log_reg_l2.pkl'.format('BOW')
        log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'BOW', model_

        print('Retraining Vectorizer with Dx_train')
        vectorizer_obj = get_vectorizer(W2V_model=None, train=Dx_train, vectorizer='BOW')

        # plotting AUC ROC
        train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy


```

```
# appending the data results
prettytable_data.append(['BOW', 'LogisticRegression', 'L2', best_parameters['logistic
```

Performing Hyperparameter Tuning...

```
CV iteration : C=0.0001, solver=saga, train_score=0.8159612740447112, test_score=0.81061690059
CV iteration : C=0.0001, solver=saga, train_score=0.7929317408214105, test_score=0.78313024540
CV iteration : C=0.0001, solver=saga, train_score=0.8173374823748631, test_score=0.81930219546
C=0.0001, penalty=l2, solver="saga", train_score=0.8087434990803283, test_score=0.804349780490
CV iteration : C=0.001, solver=saga, train_score=0.8712384409375105, test_score=0.860243025533
CV iteration : C=0.001, solver=saga, train_score=0.8698574061095923, test_score=0.858890929726
CV iteration : C=0.001, solver=saga, train_score=0.8717808662792136, test_score=0.868504117312
C=0.001, penalty=l2, solver="saga", train_score=0.8709589044421054, test_score=0.8625460241911
CV iteration : C=0.01, solver=saga, train_score=0.9184452432602117, test_score=0.9019696541070
CV iteration : C=0.01, solver=saga, train_score=0.9219529198934143, test_score=0.9044865594079
CV iteration : C=0.01, solver=saga, train_score=0.9194468596921284, test_score=0.9049199283749
C=0.01, penalty=l2, solver="saga", train_score=0.9199483409485848, test_score=0.90379204729668
CV iteration : C=0.05, solver=saga, train_score=0.9332984388798955, test_score=0.9113376041256
CV iteration : C=0.05, solver=saga, train_score=0.9385647314827107, test_score=0.9128870686391
CV iteration : C=0.05, solver=saga, train_score=0.935156004809536, test_score=0.91020035024594
C=0.05, penalty=l2, solver="saga", train_score=0.935673058390714, test_score=0.911475007670245
CV iteration : C=0.1, solver=saga, train_score=0.9370112468296754, test_score=0.91279873046257
CV iteration : C=0.1, solver=saga, train_score=0.9420622519394194, test_score=0.91272076777188
CV iteration : C=0.1, solver=saga, train_score=0.9402620774466128, test_score=0.91153270822468
C=0.1, penalty=l2, solver="saga", train_score=0.9397785254052359, test_score=0.912350735486383
CV iteration : C=0.5, solver=saga, train_score=0.9402942272175564, test_score=0.91259010699177
CV iteration : C=0.5, solver=saga, train_score=0.9468901142738392, test_score=0.91121983942968
CV iteration : C=0.5, solver=saga, train_score=0.9430591941401132, test_score=0.90981424958822
C=0.5, penalty=l2, solver="saga", train_score=0.9434145118771696, test_score=0.911208065336560
CV iteration : C=1, solver=saga, train_score=0.9405709940185543, test_score=0.9123163582958926
CV iteration : C=1, solver=saga, train_score=0.9473473620467349, test_score=0.9101417854112697
CV iteration : C=1, solver=saga, train_score=0.9433365918416323, test_score=0.9094168344107726
C=1, penalty=l2, solver="saga", train_score=0.9437516493023071, test_score=0.9106249927059783
CV iteration : C=5, solver=saga, train_score=0.9407608544319473, test_score=0.9120502886319494
CV iteration : C=5, solver=saga, train_score=0.947633200526304, test_score=0.9089810205984463
CV iteration : C=5, solver=saga, train_score=0.943541214595413, test_score=0.9090575251287862
C=5, penalty=l2, solver="saga", train_score=0.9439784231845548, test_score=0.9100296114530607
CV iteration : C=10, solver=saga, train_score=0.9407854901212498, test_score=0.912016465720273
CV iteration : C=10, solver=saga, train_score=0.947660518149659, test_score=0.908812785318535
CV iteration : C=10, solver=saga, train_score=0.943565292170752, test_score=0.9090156086437771
C=10, penalty=l2, solver="saga", train_score=0.944003766813887, test_score=0.909948286560862
CV iteration : C=50, solver=saga, train_score=0.9408044688349003, test_score=0.911985104788304
CV iteration : C=50, solver=saga, train_score=0.9476796287617055, test_score=0.908674211032224
CV iteration : C=50, solver=saga, train_score=0.943583024286492, test_score=0.9089759198880412
C=50, penalty=l2, solver="saga", train_score=0.9440223739610326, test_score=0.9098784119028567
CV iteration : C=100, solver=saga, train_score=0.9408070921474436, test_score=0.91198176353013
CV iteration : C=100, solver=saga, train_score=0.9476793796202909, test_score=0.90865562894730
```




CV iteration : C=100, solver=saga, train_score=0.943586585364281, test_score=0.9089693539491308
C=100, penalty=l2, solver="saga", train_score=0.9440243523773385, test_score=0.9098689154755211
CV iteration : C=500, solver=saga, train_score=0.9408093930416855, test_score=0.9119824669529008
CV iteration : C=500, solver=saga, train_score=0.9476820029328341, test_score=0.9086399777905998
CV iteration : C=500, solver=saga, train_score=0.9435885637408306, test_score=0.9089626121368508
C=500, penalty=l2, solver="saga", train_score=0.94402665323845, test_score=0.9098616856267867
CV iteration : C=1000, solver=saga, train_score=0.9408086749281961, test_score=0.9119828186642819
CV iteration : C=1000, solver=saga, train_score=0.947682296040381, test_score=0.9086409156876281
CV iteration : C=1000, solver=saga, train_score=0.9435890619986281, test_score=0.9089646639927819
C=1000, penalty=l2, solver="saga", train_score=0.9440266776557351, test_score=0.9098627994482281
CV iteration : C=5000, solver=saga, train_score=0.9408084404421586, test_score=0.911979653261819
CV iteration : C=5000, solver=saga, train_score=0.9476819296559474, test_score=0.9086355813982819
CV iteration : C=5000, solver=saga, train_score=0.9435888421790115, test_score=0.90896032578319
C=5000, penalty=l2, solver="saga", train_score=0.9440264040923725, test_score=0.9098585201477381
CV iteration : C=10000, solver=saga, train_score=0.9408089094142333, test_score=0.911979887736819
CV iteration : C=10000, solver=saga, train_score=0.9476817098252874, test_score=0.908639039893819
CV iteration : C=10000, solver=saga, train_score=0.9435891352718337, test_score=0.908962319014819
C=10000, penalty=l2, solver="saga", train_score=0.9440265848371182, test_score=0.909860415548019




Initializing Vectorizer
Training Model...
Saving Trained Model...
Retraining Vectorizer with Dx_train

Area Under the Curve for Train : 0.9362955972935451

Area Under the Curve for Test : 0.9125718703293875





7.2 [5.2] Logistic Regression on TFIDF, SET 2

7.2.1 [5.2.1] Applying Logistic Regression with L1 regularization on TFIDF, SET 2

```
In [53]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/TFIDF_log_reg_results_l1.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           penalty=['l1'], results_path=csv_path, ret=

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l1.pkl'.format('TFIDF')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF', mode=

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=None, train=Dx_train, vectorizer='TFIDF')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy,
```

```
# appending the data results
prettytable_data.append(['TFIDF', 'LogisticRegression', 'L1', best_parameters['logist
```

Performing Hyperparameter Tuning...

```
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
C=0.0001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
C=0.001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
CV iteration : C=0.01, solver=saga, train_score=0.6665617195095608, test_score=0.66818586446141
CV iteration : C=0.01, solver=saga, train_score=0.6685603685768778, test_score=0.66264676183429
CV iteration : C=0.01, solver=saga, train_score=0.6652055427721493, test_score=0.67068210959861
C=0.01, penalty=l1, solver="saga", train_score=0.6667758769528627, test_score=0.66717157863146
CV iteration : C=0.05, solver=saga, train_score=0.8041778010895981, test_score=0.80177124199021
CV iteration : C=0.05, solver=saga, train_score=0.8037825822012787, test_score=0.79272789227220
CV iteration : C=0.05, solver=saga, train_score=0.8036081719437248, test_score=0.81190396517050
C=0.05, penalty=l1, solver="saga", train_score=0.8038561850782006, test_score=0.80213436647764
CV iteration : C=0.1, solver=saga, train_score=0.8679350602643772, test_score=0.86422879482224
CV iteration : C=0.1, solver=saga, train_score=0.869509099084127, test_score=0.86148573909979
CV iteration : C=0.1, solver=saga, train_score=0.873027924946197, test_score=0.869633458805299
C=0.1, penalty=l1, solver="saga", train_score=0.8701573614315671, test_score=0.865115997575778
CV iteration : C=0.5, solver=saga, train_score=0.93207651015603, test_score=0.9207017732467453
CV iteration : C=0.5, solver=saga, train_score=0.9340002555897808, test_score=0.92004615391288
CV iteration : C=0.5, solver=saga, train_score=0.9337349955526095, test_score=0.91859985690943
C=0.5, penalty=l1, solver="saga", train_score=0.9332705870994734, test_score=0.919782594689688
CV iteration : C=1, solver=saga, train_score=0.9422949280377007, test_score=0.9272606630392717
CV iteration : C=1, solver=saga, train_score=0.9450826812423188, test_score=0.9251741352384211
CV iteration : C=1, solver=saga, train_score=0.9442919059134525, test_score=0.9246334271480938
C=1, penalty=l1, solver="saga", train_score=0.9438898383978239, test_score=0.9256894084752622
CV iteration : C=5, solver=saga, train_score=0.948529721251792, test_score=0.926619053542314
CV iteration : C=5, solver=saga, train_score=0.9518872608718718, test_score=0.9226894701479684
CV iteration : C=5, solver=saga, train_score=0.9507355003170677, test_score=0.9230398679120688
C=5, penalty=l1, solver="saga", train_score=0.9503841608135772, test_score=0.9241161305341171
CV iteration : C=10, solver=saga, train_score=0.9488839417220125, test_score=0.925450961408937
CV iteration : C=10, solver=saga, train_score=0.9522993554271939, test_score=0.921130333571285
CV iteration : C=10, solver=saga, train_score=0.9510467648941859, test_score=0.921627370303946
C=10, penalty=l1, solver="saga", train_score=0.9507433540144641, test_score=0.9227362217613897
CV iteration : C=50, solver=saga, train_score=0.9489942527472237, test_score=0.924226360979450
CV iteration : C=50, solver=saga, train_score=0.9524707060990112, test_score=0.918947231685538
CV iteration : C=50, solver=saga, train_score=0.9511242146724377, test_score=0.920193416145456
C=50, penalty=l1, solver="saga", train_score=0.9508630578395575, test_score=0.9211223362701485
CV iteration : C=100, solver=saga, train_score=0.9489952786236373, test_score=0.92406334275177
CV iteration : C=100, solver=saga, train_score=0.9524842037015381, test_score=0.91848748628472
CV iteration : C=100, solver=saga, train_score=0.9511243612188487, test_score=0.92000136243232
```

C=100, penalty=l1, solver="saga", train_score=0.9508679478480081, test_score=0.9208507304896063
 CV iteration : C=500, solver=saga, train_score=0.9489961872570323, test_score=0.923930864796163
 CV iteration : C=500, solver=saga, train_score=0.9524897580895484, test_score=0.91808981794363
 CV iteration : C=500, solver=saga, train_score=0.9511237896878456, test_score=0.91985521166711
 C=500, penalty=l1, solver="saga", train_score=0.9508699116781422, test_score=0.920625298135636
 CV iteration : C=1000, solver=saga, train_score=0.9489954544881652, test_score=0.9239133964639
 CV iteration : C=1000, solver=saga, train_score=0.9524900218863405, test_score=0.9180397576896
 CV iteration : C=1000, solver=saga, train_score=0.9511221776773238, test_score=0.9198372725840
 C=1000, penalty=l1, solver="saga", train_score=0.9508692180172765, test_score=0.92059680891253
 CV iteration : C=5000, solver=saga, train_score=0.9489940182611863, test_score=0.9238993280085
 CV iteration : C=5000, solver=saga, train_score=0.9524907839659621, test_score=0.9179970833747
 CV iteration : C=5000, solver=saga, train_score=0.9511221630226827, test_score=0.9198221474747
 C=5000, penalty=l1, solver="saga", train_score=0.9508689884166103, test_score=0.92057285295264
 CV iteration : C=10000, solver=saga, train_score=0.9489939449842996, test_score=0.923897569451
 CV iteration : C=10000, solver=saga, train_score=0.952490476203038, test_score=0.9179908698068
 CV iteration : C=10000, solver=saga, train_score=0.9511216501102439, test_score=0.919820212867
 C=10000, penalty=l1, solver="saga", train_score=0.9508686904325271, test_score=0.9205695507087



Initializing Vectorizer



Training Model...


Saving Trained Model...

Retraining Vectorizer with Dx_train

Area Under the Curve for Train : 0.9424295841754171

Area Under the Curve for Test : 0.9270139242219235





7.2.2 [5.2.2] Applying Logistic Regression with L2 regularization on TFIDF, SET 2

```
In [54]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/TFIDF_log_reg_results_l2.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           penalty=['l2'], results_path=csv_path, ret=

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l2.pkl'.format('TFIDF')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF', mode=

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=None, train=Dx_train, vectorizer='TFIDF')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy

# appending the data results
prettytable_data.append(['TFIDF', 'LogisticRegression', 'L2', best_parameters['logist
```

Performing Hyperparameter Tuning...

```
CV iteration : C=0.0001, solver=saga, train_score=0.8798079749873743, test_score=0.872842148060
CV iteration : C=0.0001, solver=saga, train_score=0.8810617498461918, test_score=0.871754363360
CV iteration : C=0.0001, solver=saga, train_score=0.8821356891391876, test_score=0.878267199470
C=0.0001, penalty=l2, solver="saga", train_score=0.8810018046575845, test_score=0.874287903632
CV iteration : C=0.001, solver=saga, train_score=0.8806614308865595, test_score=0.873641998372
CV iteration : C=0.001, solver=saga, train_score=0.8819202618505578, test_score=0.872573968128
CV iteration : C=0.001, solver=saga, train_score=0.8830282300558249, test_score=0.879060388347
C=0.001, penalty=l2, solver="saga", train_score=0.8818699742643141, test_score=0.8750921182827
CV iteration : C=0.01, solver=saga, train_score=0.8885272796659278, test_score=0.8812111791229
CV iteration : C=0.01, solver=saga, train_score=0.8899809392162362, test_score=0.8802692960297
CV iteration : C=0.01, solver=saga, train_score=0.8910345001319504, test_score=0.8860688832651
C=0.01, penalty=l2, solver="saga", train_score=0.8898475730047047, test_score=0.88251645280593
CV iteration : C=0.05, solver=saga, train_score=0.9084322278195993, test_score=0.8999819572058
CV iteration : C=0.05, solver=saga, train_score=0.9101220148095519, test_score=0.8990108820675
CV iteration : C=0.05, solver=saga, train_score=0.9108941766676777, test_score=0.9026329297782
C=0.05, penalty=l2, solver="saga", train_score=0.9098161397656096, test_score=0.90054192301720
CV iteration : C=0.1, solver=saga, train_score=0.9194533500287392, test_score=0.91003955463490
CV iteration : C=0.1, solver=saga, train_score=0.9213040237510908, test_score=0.90891730221891
CV iteration : C=0.1, solver=saga, train_score=0.9217368964940353, test_score=0.91107678584158
C=0.1, penalty=l2, solver="saga", train_score=0.9208314234246218, test_score=0.910011214231800
CV iteration : C=0.5, solver=saga, train_score=0.9381807972349406, test_score=0.92490909139935
CV iteration : C=0.5, solver=saga, train_score=0.9406171438017866, test_score=0.92294114895434
CV iteration : C=0.5, solver=saga, train_score=0.940288338563477, test_score=0.923392552630690
C=0.5, penalty=l2, solver="saga", train_score=0.9396954265334014, test_score=0.923747597661463
CV iteration : C=1, solver=saga, train_score=0.9426562417105522, test_score=0.9271182492369914
CV iteration : C=1, solver=saga, train_score=0.9453269424163844, test_score=0.9247991229724865
CV iteration : C=1, solver=saga, train_score=0.944764261632942, test_score=0.9249878999125791
C=1, penalty=l2, solver="saga", train_score=0.9442491485866262, test_score=0.9256350907073522
CV iteration : C=5, solver=saga, train_score=0.9476390846603046, test_score=0.9269605652986774
CV iteration : C=5, solver=saga, train_score=0.950794072310218, test_score=0.9240444089554629
CV iteration : C=5, solver=saga, train_score=0.9498137527007039, test_score=0.9241161542733007
C=5, penalty=l2, solver="saga", train_score=0.9494156365570755, test_score=0.9250403761758137
CV iteration : C=10, solver=saga, train_score=0.9483883847927599, test_score=0.926076597347087
CV iteration : C=10, solver=saga, train_score=0.9516691448909742, test_score=0.922922273776598
CV iteration : C=10, solver=saga, train_score=0.9505813921111839, test_score=0.922969401317689
C=10, penalty=l2, solver="saga", train_score=0.9502129739316393, test_score=0.9239894241471255
CV iteration : C=50, solver=saga, train_score=0.9489332863775044, test_score=0.924566583127396
CV iteration : C=50, solver=saga, train_score=0.9523606588705922, test_score=0.920663817726277
CV iteration : C=50, solver=saga, train_score=0.9510810714090181, test_score=0.920896088857789
C=50, penalty=l2, solver="saga", train_score=0.9507916722190383, test_score=0.9220421632371544
CV iteration : C=100, solver=saga, train_score=0.9489697489563174, test_score=0.92425021873517
CV iteration : C=100, solver=saga, train_score=0.9524236623377612, test_score=0.92005013997527
CV iteration : C=100, solver=saga, train_score=0.9511069661598545, test_score=0.92042926232614
C=100, penalty=l2, solver="saga", train_score=0.950833459151311, test_score=0.9215765403455288
CV iteration : C=500, solver=saga, train_score=0.9489920251298686, test_score=0.92397154607987
CV iteration : C=500, solver=saga, train_score=0.9524771544650392, test_score=0.91908082339387
```


CV iteration : C=500, solver=saga, train_score=0.9511223535330171, test_score=0.91996618775958
C=500, penalty=l2, solver="saga", train_score=0.950863844375975, test_score=0.9210061857444461
CV iteration : C=1000, solver=saga, train_score=0.9489951027591093, test_score=0.9239363749412
CV iteration : C=1000, solver=saga, train_score=0.952486387352761, test_score=0.91871656763451
CV iteration : C=1000, solver=saga, train_score=0.9511232767754069, test_score=0.9198973626499
C=1000, penalty=l2, solver="saga", train_score=0.9508682556290924, test_score=0.92085010174188
CV iteration : C=5000, solver=saga, train_score=0.9489947510300534, test_score=0.9239054243392
CV iteration : C=5000, solver=saga, train_score=0.9524904029261512, test_score=0.9181687771499
CV iteration : C=5000, solver=saga, train_score=0.9511224854247872, test_score=0.9198368035883
C=5000, penalty=l2, solver="saga", train_score=0.9508692131269972, test_score=0.92063700169250
CV iteration : C=10000, solver=saga, train_score=0.9489939449842997, test_score=0.923899562482
CV iteration : C=10000, solver=saga, train_score=0.9524901391293592, test_score=0.918076101199
CV iteration : C=10000, solver=saga, train_score=0.9511222949144527, test_score=0.919826837431
C=10000, penalty=l2, solver="saga", train_score=0.9508687930093705, test_score=0.9206008337044



Initializing Vectorizer


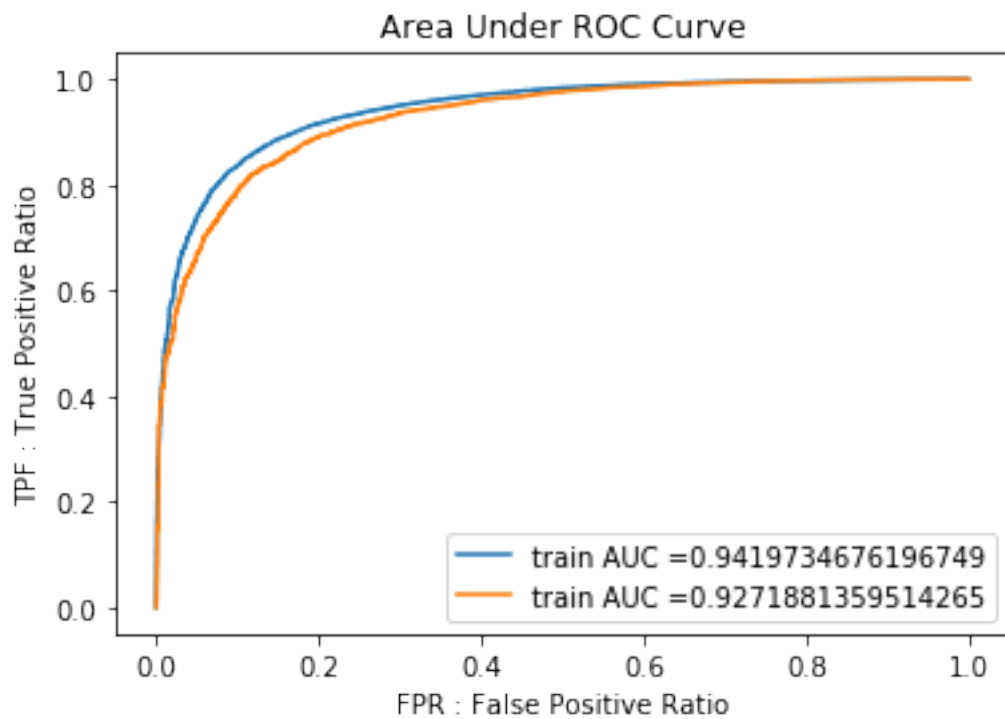
Training Model...


Saving Trained Model...

Retraining Vectorizer with Dx_train

Area Under the Curve for Train : 0.9419734676196749

Area Under the Curve for Test : 0.9271881359514265





7.2.3 [5.2.3] Feature Importance on TFIDF, SET 2

```
In [55]: # Please write all the code with proper documentation
feature_names = vectorizer_obj.get_feature_names()
weights = np.reshape(log_reg.coef_,(log_reg.coef_.shape[1], log_reg.coef_.shape[0]))

# making a list of feature names along with their feature weights
features_with_weights = [(feature_names[i],weights[i]) for i in range(len(feature_names))]
features_with_weights.sort(key=lambda x : abs(x[1][0]), reverse=True)
```

[5.2.3.1] Top 10 important features of positive class from SET 2

```
In [56]: positive_weights = [i for i in features_with_weights if i[1][0]>=0]
print('Top 10 features of positive class with the feature names : ')
for i in positive_weights[:10]:
    print(i[0], ' : ',i[1][0])
```

Top 10 features of positive class with the feature names :

```
great : 8.250425983875289
best : 7.047466209780987
delicious : 6.5610791490084805
excellent : 5.676299237656143
loves : 5.533156858346458
love : 5.46371508484913
```

```
perfect : 5.178975977470662
good : 4.664923465749984
wonderful : 4.549403455960735
pleased : 4.0260402569712666
```

[5.2.3.2] Top 10 important features of negative class from SET 2

```
In [57]: # Please write all the code with proper documentation
negative_weights = [i for i in features_with_weights if i[1][0]<0]
print('Top 10 features of negative class with the feature names : ')
for i in negative_weights[:10]:
    print(i[0], ' : ', i[1][0])
```

```
Top 10 features of negative class with the feature names :
not : -5.3249053217795
disappointed : -5.027833945426516
awful : -4.979721775746406
terrible : -4.863041733252396
money : -4.474683323856104
unfortunately : -4.223315676440612
not buy : -4.206004932916333
not good : -4.107539907827749
horrible : -4.097944215892398
bland : -3.255097383219475
```

7.3 Preparing/Training Google Word2Vec

```
In [58]: is_your_ram_gt_16g=True
want_to_use_google_w2v = False
want_to_train_w2v = True

path_to_word2vec = '/home/monodeepdas112/Datasets/GoogleNews-vectors-negative300.bin.'

if want_to_train_w2v:

    # Train your own Word2Vec model using your own text corpus
    i=0
    list_of_sentences=[]
    for sentence in preprocessed_reviews:
        list_of_sentences.append(sentence.split())

    # min_count = 5 considers only words that occurred atleast 5 times
    w2v_model=Word2Vec(list_of_sentences,min_count=5,size=300, workers=4)
    print(w2v_model.wv.most_similar('great'))
    print('='*50)
    print(w2v_model.wv.most_similar('worst'))
```

```

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile(path_to_word2vec):
        print('Preparing to load pre-trained Word2Vec model !')
        w2v_model=KeyedVectors.load_word2vec_format(path_to_word2vec, binary=True)
        print('Successfully loaded model into memory !!!')
        print('Words similar to "similar" : ', w2v_model.wv.most_similar('great'))
        print('Words similar to "worst" : ',w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True, ")

[('terrific', 0.78861403465271), ('fantastic', 0.7743387222290039), ('excellent', 0.7511157989)
=====
[('nastiest', 0.7911990284919739), ('greatest', 0.6833856105804443), ('disgusting', 0.63364648)

```

7.4 [5.3] Logistic Regression on AVG W2V, SET 3

7.4.1 [5.3.1] Applying Logistic Regression with L1 regularization on AVG W2V SET 3

```

In [59]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/Avg-W2Vec_log_reg_results_l1.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='Avg-W2Vec',
                                           penalty=['l1'], results_path=csv_path, ret=

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l1.pkl'.format('Avg-W2Vec')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'Avg-W2Vec', r

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=w2v_model, train=Dx_train, vectorizer='Avg-W

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy

# appending the data results
prettytable_data.append(['Avg Word2Vec', 'LogisticRegression', 'L1', best_parameters[

```

Performing Hyperparameter Tuning...


```

CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.0001, solver=saga, train_score=0.5, test_score=0.5
C=0.0001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5
CV iteration : C=0.001, solver=saga, train_score=0.5, test_score=0.5



```


C=0.001, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
 CV iteration : C=0.01, solver=saga, train_score=0.6014673916388728, test_score=0.6049558948059
 CV iteration : C=0.01, solver=saga, train_score=0.6142678012273586, test_score=0.6107555569519
 CV iteration : C=0.01, solver=saga, train_score=0.616261173724205, test_score=0.61250774429044
 C=0.01, penalty=l1, solver="saga", train_score=0.6106654555301455, test_score=0.60940639868278
 CV iteration : C=0.05, solver=saga, train_score=0.6390860232545665, test_score=0.6371251181603
 CV iteration : C=0.05, solver=saga, train_score=0.6383272117822198, test_score=0.6327745069270
 CV iteration : C=0.05, solver=saga, train_score=0.6381474845425776, test_score=0.6412530730939
 C=0.05, penalty=l1, solver="saga", train_score=0.638520239859788, test_score=0.637050899393764
 CV iteration : C=0.1, solver=saga, train_score=0.6533847766388596, test_score=0.64974569508193
 CV iteration : C=0.1, solver=saga, train_score=0.6500071225133854, test_score=0.64770354153435
 CV iteration : C=0.1, solver=saga, train_score=0.6513110305718088, test_score=0.65119683582713
 C=0.1, penalty=l1, solver="saga", train_score=0.6515676432413513, test_score=0.649548690814478
 CV iteration : C=0.5, solver=saga, train_score=0.6621460689734817, test_score=0.65838278883916
 CV iteration : C=0.5, solver=saga, train_score=0.6614628206266815, test_score=0.65880982508103
 CV iteration : C=0.5, solver=saga, train_score=0.6599744001005425, test_score=0.66042778381739
 C=0.5, penalty=l1, solver="saga", train_score=0.6611944299002352, test_score=0.659206799245865
 CV iteration : C=1, solver=saga, train_score=0.6633257535721749, test_score=0.659254798603612
 CV iteration : C=1, solver=saga, train_score=0.6628535573143962, test_score=0.659396479673837
 CV iteration : C=1, solver=saga, train_score=0.6606331555275957, test_score=0.6616021488911067
 C=1, penalty=l1, solver="saga", train_score=0.6622708221380557, test_score=0.660084475722852
 CV iteration : C=5, solver=saga, train_score=0.663963995255175, test_score=0.6596195819300538
 CV iteration : C=5, solver=saga, train_score=0.6637815358073369, test_score=0.6594682874152624
 CV iteration : C=5, solver=saga, train_score=0.6613421763736469, test_score=0.6631163013034327
 C=5, penalty=l1, solver="saga", train_score=0.663029235812053, test_score=0.660734723549583
 CV iteration : C=10, solver=saga, train_score=0.6639943905077706, test_score=0.659610144341180
 CV iteration : C=10, solver=saga, train_score=0.6638383986714022, test_score=0.659423737306296
 CV iteration : C=10, solver=saga, train_score=0.661417750357837, test_score=0.6632802152783771
 C=10, penalty=l1, solver="saga", train_score=0.6630835131790033, test_score=0.6607713656419514
 CV iteration : C=50, solver=saga, train_score=0.6640088113990698, test_score=0.659589510606501
 CV iteration : C=50, solver=saga, train_score=0.6638804596043576, test_score=0.659398414086463
 CV iteration : C=50, solver=saga, train_score=0.6614706536122342, test_score=0.663352206108574
 C=50, penalty=l1, solver="saga", train_score=0.6631199748718872, test_score=0.660780043600513
 CV iteration : C=100, solver=saga, train_score=0.6640104381459541, test_score=0.65957555938816
 CV iteration : C=100, solver=saga, train_score=0.6638910554421718, test_score=0.65939935198349
 CV iteration : C=100, solver=saga, train_score=0.6614777904224534, test_score=0.66337506964585
 C=100, penalty=l1, solver="saga", train_score=0.6631264280035264, test_score=0.660783327005837
 CV iteration : C=500, solver=saga, train_score=0.6640097493432193, test_score=0.65957221812999
 CV iteration : C=500, solver=saga, train_score=0.6638894433506648, test_score=0.65939771066368
 CV iteration : C=500, solver=saga, train_score=0.6614800618918251, test_score=0.66338573929658
 C=500, penalty=l1, solver="saga", train_score=0.6631264181952364, test_score=0.660785222696755
 CV iteration : C=1000, solver=saga, train_score=0.66401057004435, test_score=0.659572159511430
 CV iteration : C=1000, solver=saga, train_score=0.663888710581798, test_score=0.65939501420972
 CV iteration : C=1000, solver=saga, train_score=0.6614787576287666, test_score=0.6633857392965
 C=1000, penalty=l1, solver="saga", train_score=0.6631260127516382, test_score=0.66078430433924
 CV iteration : C=5000, solver=saga, train_score=0.6640112002255756, test_score=0.6595725698413
 CV iteration : C=5000, solver=saga, train_score=0.6638894726614195, test_score=0.6593946624983
 CV iteration : C=5000, solver=saga, train_score=0.6614792119226409, test_score=0.6633884360214

C=5000, penalty=l1, solver="saga", train_score=0.6631266282698787, test_score=0.66078522278707
CV iteration : C=10000, solver=saga, train_score=0.6640112002255756, test_score=0.659572745697
CV iteration : C=10000, solver=saga, train_score=0.6638902933625503, test_score=0.659394545261
CV iteration : C=10000, solver=saga, train_score=0.6614793438144109, test_score=0.663388436021
C=10000, penalty=l1, solver="saga", train_score=0.6631269458008456, test_score=0.6607852423265



Initializing Vectorizer
Training Model...
Saving Trained Model...
Retraining Vectorizer with Dx_train
Area Under the Curve for Train : 0.6627319168112705
Area Under the Curve for Test : 0.6506057032390009





7.4.2 [5.3.2] Applying Logistic Regression with L2 regularization on AVG W2V, SET 3

```
In [60]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/Avg-W2Vec_log_reg_results_l2.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='Avg-W2Vec',
                                           penalty=['l2'], results_path=csv_path, ret=1)

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l2.pkl'.format('Avg-W2Vec')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'Avg-W2Vec', n_iter=1000)

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=w2v_model, train=Dx_train, vectorizer='Avg-W2Vec')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy_test)

# appending the data results
prettytable_data.append(['Avg-Word2Vec', 'LogisticRegression', 'L2', best_parameters['C'], test_score])
```

Performing Hyperparameter Tuning...



CV iteration : C=0.0001, solver=saga, train_score=0.6147310283942073, test_score=0.616508500107
CV iteration : C=0.0001, solver=saga, train_score=0.616703715460573, test_score=0.6124426578548
CV iteration : C=0.0001, solver=saga, train_score=0.6162938535738742, test_score=0.614984217124
C=0.0001, penalty=l2, solver="saga", train_score=0.6159095324762182, test_score=0.614645125027
CV iteration : C=0.001, solver=saga, train_score=0.6243115270111794, test_score=0.624784562121
CV iteration : C=0.001, solver=saga, train_score=0.6253406422631185, test_score=0.621191595832
CV iteration : C=0.001, solver=saga, train_score=0.6254126600389391, test_score=0.623550481048
C=0.001, penalty=l2, solver="saga", train_score=0.625021609771079, test_score=0.62317554633429
CV iteration : C=0.01, solver=saga, train_score=0.6490087835538528, test_score=0.6467836404027
CV iteration : C=0.01, solver=saga, train_score=0.6473274453889347, test_score=0.6450583788145
CV iteration : C=0.01, solver=saga, train_score=0.6479435699010554, test_score=0.6461223030405
C=0.01, penalty=l2, solver="saga", train_score=0.6480932662812809, test_score=0.64598810741928
CV iteration : C=0.05, solver=saga, train_score=0.6593885717953746, test_score=0.6557646492775
CV iteration : C=0.05, solver=saga, train_score=0.6576928273945348, test_score=0.6562435043303
CV iteration : C=0.05, solver=saga, train_score=0.657269871429558, test_score=0.65675806883542
C=0.05, penalty=l2, solver="saga", train_score=0.6581170902064891, test_score=0.65625540748108
CV iteration : C=0.1, solver=saga, train_score=0.661128560780393, test_score=0.657414586010492
CV iteration : C=0.1, solver=saga, train_score=0.6599546943665022, test_score=0.65794901146239
CV iteration : C=0.1, solver=saga, train_score=0.6589422737273587, test_score=0.65903521852382
C=0.1, penalty=l2, solver="saga", train_score=0.6600085096247513, test_score=0.65813293866557
CV iteration : C=0.5, solver=saga, train_score=0.663154153758943, test_score=0.659198583400324
CV iteration : C=0.5, solver=saga, train_score=0.6627829623617668, test_score=0.65924811608726
CV iteration : C=0.5, solver=saga, train_score=0.660721332503138, test_score=0.661865958936618
C=0.5, penalty=l2, solver="saga", train_score=0.662219482874616, test_score=0.6601042194747365
CV iteration : C=1, solver=saga, train_score=0.6636040591877931, test_score=0.6595050412551595
CV iteration : C=1, solver=saga, train_score=0.6633454210885604, test_score=0.6593875110334793
CV iteration : C=1, solver=saga, train_score=0.6610751981219549, test_score=0.6625516305571293
C=1, penalty=l2, solver="saga", train_score=0.6626748927994361, test_score=0.6604813942819227
CV iteration : C=5, solver=saga, train_score=0.6639661642510208, test_score=0.659624681745159
CV iteration : C=5, solver=saga, train_score=0.6638100405162561, test_score=0.6594032794273108
CV iteration : C=5, solver=saga, train_score=0.6613922219730282, test_score=0.6632366573086403
C=5, penalty=l2, solver="saga", train_score=0.6630561422467683, test_score=0.6607548728270367
CV iteration : C=10, solver=saga, train_score=0.664000633698516, test_score=0.6596116684238558
CV iteration : C=10, solver=saga, train_score=0.6638514566126088, test_score=0.659405155221372
CV iteration : C=10, solver=saga, train_score=0.6614448321346031, test_score=0.663310582745838
C=10, penalty=l2, solver="saga", train_score=0.663098974148576, test_score=0.660775802130355
CV iteration : C=50, solver=saga, train_score=0.6640082691501084, test_score=0.659583941842880
CV iteration : C=50, solver=saga, train_score=0.663887362287083, test_score=0.6593909109102162
CV iteration : C=50, solver=saga, train_score=0.6614730716280169, test_score=0.663370848685123
C=50, penalty=l2, solver="saga", train_score=0.663122901021736, test_score=0.660781900479407
CV iteration : C=100, solver=saga, train_score=0.664009734687842, test_score=0.659579545450548
CV iteration : C=100, solver=saga, train_score=0.6638882709204779, test_score=0.65939319703422
CV iteration : C=100, solver=saga, train_score=0.6614753870613118, test_score=0.66337213842312
C=100, penalty=l2, solver="saga", train_score=0.6631244642232105, test_score=0.660781626969300
CV iteration : C=500, solver=saga, train_score=0.6640081812178442, test_score=0.65958142124461
CV iteration : C=500, solver=saga, train_score=0.6638886666156659, test_score=0.65939495559116
CV iteration : C=500, solver=saga, train_score=0.6614789041751776, test_score=0.66337905610876
C=500, penalty=l2, solver="saga", train_score=0.6631252506695625, test_score=0.660785144314844


CV iteration : C=1000, solver=saga, train_score=0.6640073605167134, test_score=0.6595801316361
 CV iteration : C=1000, solver=saga, train_score=0.6638887252371752, test_score=0.6593935487456
 CV iteration : C=1000, solver=saga, train_score=0.6614791826133586, test_score=0.6633802872223
 C=1000, penalty=l2, solver="saga", train_score=0.6631250894557491, test_score=0.66078465586803
 CV iteration : C=5000, solver=saga, train_score=0.6640075803473735, test_score=0.6595824177602
 CV iteration : C=5000, solver=saga, train_score=0.6638895019721742, test_score=0.6593917315701
 CV iteration : C=5000, solver=saga, train_score=0.6614785964277143, test_score=0.6633799941000
 C=5000, penalty=l2, solver="saga", train_score=0.6631252262490873, test_score=0.66078471447678
 CV iteration : C=10000, solver=saga, train_score=0.6640077122457695, test_score=0.659582241904
 CV iteration : C=10000, solver=saga, train_score=0.6638896631813247, test_score=0.659392024662
 CV iteration : C=10000, solver=saga, train_score=0.6614784938452265, test_score=0.663379876851
 C=10000, penalty=l2, solver="saga", train_score=0.6631252897574402, test_score=0.6607847144728



```

Initializing Vectorizer
Training Model...
Saving Trained Model...
Retraining Vectorizer with Dx_train
Area Under the Curve for Train :  0.6627308160480285
Area Under the Curve for Test :  0.6506013837971308
  
```





7.5 [5.4] Logistic Regression on TFIDF W2V, SET 4

7.5.1 [5.4.1] Applying Logistic Regression with L1 regularization on TFIDF W2V, SET 4

```
In [61]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/TFIDF-W2Vec_log_reg_results_l1.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF-W2Vec',
                                           penalty=['l1'], results_path=csv_path, return_results=True)

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_l1.pkl'.format('TFIDF-W2Vec')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF-W2Vec')


print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=w2v_model, train=Dx_train, vectorizer='TFIDF-W2Vec')

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy_train)



# appending the data results
prettytable_data.append(['TFIDF-Word2Vec', 'LogisticRegression', 'L1', best_parameters['penalty'], train_score, test_score])
```


Performing Hyperparameter Tuning...

CV iteration : C=500, solver=saga, train_score=0.5, test_score=0.5
 C=500, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
 CV iteration : C=1000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=1000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=1000, solver=saga, train_score=0.5, test_score=0.5
 C=1000, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
 CV iteration : C=5000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=5000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=5000, solver=saga, train_score=0.5, test_score=0.5
 C=5000, penalty=l1, solver="saga", train_score=0.5, test_score=0.5
 CV iteration : C=10000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=10000, solver=saga, train_score=0.5, test_score=0.5
 CV iteration : C=10000, solver=saga, train_score=0.5, test_score=0.5
 C=10000, penalty=l1, solver="saga", train_score=0.5, test_score=0.5



Initializing Vectorizer
 Training Model...
 Saving Trained Model...
 Retraining Vectorizer with Dx_train
 Area Under the Curve for Train : 0.5
 Area Under the Curve for Test : 0.5





7.5.2 [5.4.2] Applying Logistic Regression with L2 regularization on TFIDF W2V, SET 4

```
In [62]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment5/Avg-W2Vec_log_reg_results_12.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='Avg-W2Vec',
                                           penalty=['l2'], results_path=csv_path, ret=1)

# Analysing best parameters
best_parameters = analyse_results(cv_results)

# retraining the model with best parameters
model_path = 'saved_models/Assignment5/{0}_log_reg_12.pkl'.format('Avg-W2Vec')
log_reg = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'Avg-W2Vec', n=1000)

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model=w2v_model, train=Dx_train, vectorizer='Avg-W2Vec')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(log_reg, vectorizer_obj, Dx_train, Dx_test, Dy_test)

# appending the data results
prettytable_data.append(['TFIDF-Word2Vec', 'LogisticRegression', 'L1', best_parameters['L1'],
                        test_score])
```

Performing Hyperparameter Tuning...

CV iteration : C=0.0001, solver=saga, train_score=0.6147316439200554, test_score=0.61651043451
 CV iteration : C=0.0001, solver=saga, train_score=0.6167041551218931, test_score=0.61244177857
 CV iteration : C=0.0001, solver=saga, train_score=0.6162951285276506, test_score=0.61498521374
 C=0.0001, penalty=l2, solver="saga", train_score=0.6159103091898663, test_score=0.614645808943
 CV iteration : C=0.001, solver=saga, train_score=0.624311336491274, test_score=0.6247843276468
 CV iteration : C=0.001, solver=saga, train_score=0.6253401732910437, test_score=0.621191419977
 CV iteration : C=0.001, solver=saga, train_score=0.6254125721110925, test_score=0.623550656922
 C=0.001, penalty=l2, solver="saga", train_score=0.6250213606311368, test_score=0.6231754681820
 CV iteration : C=0.01, solver=saga, train_score=0.6490099120179077, test_score=0.6467838162584
 CV iteration : C=0.01, solver=saga, train_score=0.6473275186658213, test_score=0.6450580271031
 CV iteration : C=0.01, solver=saga, train_score=0.647943379390721, test_score=0.64612236166504
 C=0.01, penalty=l2, solver="saga", train_score=0.64809360335815, test_score=0.6459880683422098
 CV iteration : C=0.05, solver=saga, train_score=0.6593886450722612, test_score=0.6557645320403
 CV iteration : C=0.05, solver=saga, train_score=0.6576939558585897, test_score=0.6562437974231
 CV iteration : C=0.05, solver=saga, train_score=0.6572695929913768, test_score=0.6567597689446
 C=0.05, penalty=l2, solver="saga", train_score=0.6581173979740759, test_score=0.65625603280271
 CV iteration : C=0.1, solver=saga, train_score=0.6611285754357703, test_score=0.65741446877336
 CV iteration : C=0.1, solver=saga, train_score=0.6599532727949007, test_score=0.65795288028764
 CV iteration : C=0.1, solver=saga, train_score=0.6589434754079295, test_score=0.65903697725746
 C=0.1, penalty=l2, solver="saga", train_score=0.6600084412128668, test_score=0.658134775439490
 CV iteration : C=0.5, solver=saga, train_score=0.6631543589342257, test_score=0.65919922820453
 CV iteration : C=0.5, solver=saga, train_score=0.6627829623617668, test_score=0.65924799885013
 CV iteration : C=0.5, solver=saga, train_score=0.6607212445752915, test_score=0.66186613480998
 C=0.5, penalty=l2, solver="saga", train_score=0.6622195219570947, test_score=0.660104453954884
 CV iteration : C=1, solver=saga, train_score=0.6636037367694918, test_score=0.6595052757294172
 CV iteration : C=1, solver=saga, train_score=0.663345406433183, test_score=0.6593882730748168
 CV iteration : C=1, solver=saga, train_score=0.6610757696529581, test_score=0.6625511029370383
 C=1, penalty=l2, solver="saga", train_score=0.6626749709518777, test_score=0.6604815505804241
 CV iteration : C=5, solver=saga, train_score=0.6639657978665874, test_score=0.6596247403637234
 CV iteration : C=5, solver=saga, train_score=0.6638098646517281, test_score=0.6594037483758262
 CV iteration : C=5, solver=saga, train_score=0.6613922512823104, test_score=0.663235250321731
 C=5, penalty=l2, solver="saga", train_score=0.6630559712668753, test_score=0.6607545796870936
 CV iteration : C=10, solver=saga, train_score=0.6640005164554973, test_score=0.659611668423855
 CV iteration : C=10, solver=saga, train_score=0.6638511341943074, test_score=0.659404393180034
 CV iteration : C=10, solver=saga, train_score=0.6614454036656063, test_score=0.663310641370292
 C=10, penalty=l2, solver="saga", train_score=0.6630990181051369, test_score=0.6607755676580611
 CV iteration : C=50, solver=saga, train_score=0.6640082838054856, test_score=0.659583824605752
 CV iteration : C=50, solver=saga, train_score=0.6638883735081191, test_score=0.659390852291651
 CV iteration : C=50, solver=saga, train_score=0.661472954390888, test_score=0.6633709073095784
 C=50, penalty=l2, solver="saga", train_score=0.6631232039014976, test_score=0.6607818614023274
 CV iteration : C=100, solver=saga, train_score=0.6640098958969927, test_score=0.65957960406911
 CV iteration : C=100, solver=saga, train_score=0.6638891795538727, test_score=0.65939243499289
 CV iteration : C=100, solver=saga, train_score=0.6614752844788241, test_score=0.66337213842312
 C=100, penalty=l2, solver="saga", train_score=0.6631247866432298, test_score=0.6607813924950422
 CV iteration : C=500, solver=saga, train_score=0.6640079613871842, test_score=0.65958171433743
 CV iteration : C=500, solver=saga, train_score=0.6638886519602887, test_score=0.65939448664264
 CV iteration : C=500, solver=saga, train_score=0.661478962793742, test_score=0.663379056108762
 C=500, penalty=l2, solver="saga", train_score=0.6631251920470715, test_score=0.660785085696280

CV iteration : C=1000, solver=saga, train_score=0.6640075510366188, test_score=0.65958048334751
 CV iteration : C=1000, solver=saga, train_score=0.66388875454793, test_score=0.6593935487456151
 CV iteration : C=1000, solver=saga, train_score=0.6614794903608219, test_score=0.66338022859781
 C=1000, penalty=l2, solver="saga", train_score=0.6631252653151235, test_score=0.660784753563681
 CV iteration : C=5000, solver=saga, train_score=0.6640076243135056, test_score=0.65958235914161
 CV iteration : C=5000, solver=saga, train_score=0.6638895899044381, test_score=0.65939208328151
 CV iteration : C=5000, solver=saga, train_score=0.6614784498813032, test_score=0.66337993547551
 C=5000, penalty=l2, solver="saga", train_score=0.6631252213664157, test_score=0.660784792632901
 CV iteration : C=10000, solver=saga, train_score=0.6640076096581282, test_score=0.6595825936151
 CV iteration : C=10000, solver=saga, train_score=0.6638897804243435, test_score=0.6593919660441
 CV iteration : C=10000, solver=saga, train_score=0.6614786110823554, test_score=0.6633797009771
 C=10000, penalty=l2, solver="saga", train_score=0.6631253337216091, test_score=0.66078475354601






Loading Model...

Retraining Vectorizer with Dx_train

Area Under the Curve for Train : 0.6627308160480285

Area Under the Curve for Test : 0.6506013837971308





8 [6] Conclusions

In [63]: `from prettytable import PrettyTable`

In [64]: `# Please compare all your models using Prettytable library`
`x = PrettyTable()`

```
x.field_names = ["Vectorizer", "Model", "Penalty", "Hyper parameter: 1/C", "Train AUC"]
[x.add_row(i) for i in prettytable_data]
print(x)
```

Vectorizer	Model	Penalty	Hyper parameter: 1/C	Train AUC
BOW	LogisticRegression	L2	0.1	0.9362955972935451
TFIDF	LogisticRegression	L1	1.0	0.9424295841754171
TFIDF	LogisticRegression	L2	1.0	0.9419734676196749
Avg Word2Vec	LogisticRegression	L1	10000.0	0.6627319168112705
Avg-Word2Vec	LogisticRegression	L2	500.0	0.6627308160480285
TFIDF-Word2Vec	LogisticRegression	L1	0.0001	0.5
TFIDF-Word2Vec	LogisticRegression	L1	500.0	0.6627308160480285