

Assignment09 - Amazon Fine Food Reviews Analysis_RF

June 16, 2019

1 Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews>

EDA: <https://nycdatascience.com/blog/student-works/amazon-fine-foods-visualization/>

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

1. Id
2. ProductId - unique identifier for the product
3. UserId - unique identifier for the user
4. ProfileName
5. HelpfulnessNumerator - number of users who found the review helpful
6. HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
7. Score - rating between 1 and 5
8. Time - timestamp for the review
9. Summary - brief summary of the review
10. Text - text of the review

Objective: Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

[Q] How to determine if a review is positive or negative? [Ans] We could use Score/Rating. A rating of 4 or 5 can be considered as a positive review. A rating of 1 or 2 can be considered as negative one. A review of rating 3 is considered neutral and such reviews are ignored from our analysis. This is an approximate and proxy way of determining the polarity (positivity/negativity) of a review.

2 [1]. Reading Data

2.1 [1.1] Loading the data

The dataset is available in two forms 1. .csv file 2. SQLite Database

In order to load the data, We have used the SQLITE dataset as it is easier to query the data and visualise the data efficiently.

Here as we only want to get the global sentiment of the recommendations (positive or negative), we will purposefully ignore all Scores equal to 3. If the score is above 3, then the recommendation will be set to "positive". Otherwise, it will be set to "negative".

```
In [59]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import sqlite3
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer

from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer

import re
# Tutorial about Python regular expressions: https://pymotw.com/2/re/
import string
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

from tqdm import tqdm
import os

In [60]: # using SQLite Table to read data.
db_path = '/home/monodeepdas112/Datasets/amazon-fine-food-reviews/database.sqlite'
con = sqlite3.connect(db_path)

# filtering only positive and negative reviews i.e.
# not taking into consideration those reviews with Score=3
# SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
# you can change the number to any other number based on your computing power
```

```

# filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000
# for tsne assignment you can take 5k data points

filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 100000

# Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0)
def partition(x):
    if x < 3:
        return 0
    return 1

#changing reviews with score less than 3 to be positive and vice-versa
actualScore = filtered_data['Score']
positiveNegative = actualScore.map(partition)
filtered_data['Score'] = positiveNegative
print("Number of data points in our data", filtered_data.shape)
filtered_data.head(3)

```

Number of data points in our data (100000, 10)

```

Out[60]:
  Id  ProductId  UserId  ProfileName \
0   1  B001E4KFG0  A3SGXH7AUHU8GW  delmartian
1   2  B00813GRG4  A1D87F6ZCVE5NK  dll pa
2   3  B000LQOCHO  ABXLMWJIXXAIN  Natalia Corres "Natalia Corres"

  HelpfulnessNumerator  HelpfulnessDenominator  Score  Time \
0                      1                      1      1  1303862400
1                      0                      0      0  1346976000
2                      1                      1      1  1219017600

  Summary  Text
0  Good Quality Dog Food  I have bought several of the Vitality canned d...
1  Not as Advertised  Product arrived labeled as Jumbo Salted Peanut...
2  "Delight" says it all  This is a confection that has been around a fe...

```

```

In [61]: display = pd.read_sql_query("""
SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
FROM Reviews
GROUP BY UserId
HAVING COUNT(*)>1
""", con)

```

```

In [62]: print(display.shape)
display.head()

```

(80668, 7)

```
Out [62]:
```

	UserId	ProductId	ProfileName	Time	Score	\
0	#oc-R115TNMSPFT9I7	B005ZBZLT4	Breyton	1331510400	2	
1	#oc-R11D9D7SHXIJB9	B005HG9ESG	Louis E. Emory "hoppy"	1342396800	5	
2	#oc-R11DNU2NBKQ23Z	B005ZBZLT4	Kim Cieszykowski	1348531200	1	
3	#oc-R1105J5ZVQE25C	B005HG9ESG	Penguin Chick	1346889600	5	
4	#oc-R12KPBODL2B5ZD	B0070SBEV0	Christopher P. Presta	1348617600	1	

	Text	COUNT(*)
0	Overall its just OK when considering the price...	2
1	My wife has recurring extreme muscle spasms, u...	3
2	This coffee is horrible and unfortunately not ...	2
3	This will be the bottle that you grab from the...	3
4	I didnt like this coffee. Instead of telling y...	2

```
In [63]: display[display['UserId']=='AZY10LLTJ71NX']
```

```
Out [63]:
```

	UserId	ProductId	ProfileName	Time	\
80638	AZY10LLTJ71NX	B001ATMQK2	undertheshrine "undertheshrine"	1296691200	

	Score	Text	COUNT(*)
80638	5	I bought this 6 pack because for the price tha...	5

```
In [64]: display['COUNT(*)'].sum()
```

```
Out [64]: 393063
```

3 [2] Exploratory Data Analysis

3.1 [2.1] Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [65]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND UserId="AR5J8UI46CURR"
ORDER BY ProductID
""", con)
display.head()
```

```
Out [65]:
```

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	\
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	

	HelpfulnessDenominator	Score	Time \
0	2	5	1199577600
1	2	5	1199577600
2	2	5	1199577600
3	2	5	1199577600
4	2	5	1199577600

	Summary \
0	LOACKER QUADRATINI VANILLA WAFERS
1	LOACKER QUADRATINI VANILLA WAFERS
2	LOACKER QUADRATINI VANILLA WAFERS
3	LOACKER QUADRATINI VANILLA WAFERS
4	LOACKER QUADRATINI VANILLA WAFERS

	Text
0	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
1	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
2	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
3	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...
4	DELICIOUS WAFERS. I FIND THAT EUROPEAN WAFERS ...

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that ProductId=B000HDOPZG was Loacker Quadratini Vanilla Wafer Cookies, 8.82-Ounce Packages (Pack of 8) ProductId=B000HDL1RQ was Loacker Quadratini Lemon Wafer Cookies, 8.82-Ounce Packages (Pack of 8) and so on

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [66]: #Sorting data according to ProductId in ascending order
sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False)

In [67]: #Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first')
final.shape

Out[67]: (87775, 10)

In [68]: #Checking to see how much % of data still remains
(final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100

Out[68]: 87.775
```

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations

```
In [69]: display= pd.read_sql_query("""
SELECT *
FROM Reviews
WHERE Score != 3 AND Id=44737 OR Id=64422
ORDER BY ProductID
""", con)
```

```
display.head()
```

```
Out [69]:
```

	Id	ProductId	UserId	ProfileName	\
0	64422	B000MIDR0Q	A161DK06JJMCYF	J. E. Stephens	"Jeanne"
1	44737	B001EQ55RW	A2V0I904FH7ABY		Ram

	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	\
0	3	1	5	1224892800	
1	3	2	4	1212883200	

	Summary	\
0	Bought This for My Son at College	
1	Pure cocoa taste with crunchy almonds inside	

	Text
0	My son loves spaghetti so I didn't hesitate or...
1	It was almost a 'love at first bite' - the per...

```
In [70]: final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [71]: #Before starting the next phase of preprocessing lets see the number of entries left
print(final.shape)
```

```
#How many positive and negative reviews are present in our dataset?
final['Score'].value_counts()
```

```
(87773, 10)
```

```
Out [71]: 1    73592
0    14181
Name: Score, dtype: int64
```

4 [3] Preprocessing

4.1 [3.1]. Preprocessing Review Text

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

```
In [72]: # printing some random reviews
sent_0 = final['Text'].values[0]
print(sent_0)
print("="*50)

sent_1000 = final['Text'].values[1000]
print(sent_1000)
print("="*50)

sent_1500 = final['Text'].values[1500]
print(sent_1500)
print("="*50)

sent_4900 = final['Text'].values[4900]
print(sent_4900)
print("="*50)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
=====
The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste
=====
was way to hot for my blood, took a bite and did a jig lol
=====
My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid
=====
```

```
In [73]: # remove urls from text python: https://stackoverflow.com/a/40823105/4084039
sent_0 = re.sub(r"http\S+", "", sent_0)
sent_1000 = re.sub(r"http\S+", "", sent_1000)
sent_150 = re.sub(r"http\S+", "", sent_1500)
sent_4900 = re.sub(r"http\S+", "", sent_4900)

print(sent_0)
```

```
My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
```

```
In [74]: # https://stackoverflow.com/questions/16206380/python-beautifulsoup-how-to-remove-all
from bs4 import BeautifulSoup

soup = BeautifulSoup(sent_0, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1000, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_1500, 'lxml')
text = soup.get_text()
print(text)
print("="*50)

soup = BeautifulSoup(sent_4900, 'lxml')
text = soup.get_text()
print(text)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its
=====

The Candy Blocks were a nice visual for the Lego Birthday party but the candy has little taste
=====

was way to hot for my blood, took a bite and did a jig lol
=====

My dog LOVES these treats. They tend to have a very strong fish oil smell. So if you are afraid

```
In [75]: # https://stackoverflow.com/a/47091490/4084039
import re

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"'\re", " are", phrase)
    phrase = re.sub(r"'\s", " is", phrase)
    phrase = re.sub(r"'\d", " would", phrase)
    phrase = re.sub(r"'\ll", " will", phrase)
    phrase = re.sub(r"'\t", " not", phrase)
    phrase = re.sub(r"'\ve", " have", phrase)
    phrase = re.sub(r"'\m", " am", phrase)
    return phrase
```



```
In [76]: sent_1500 = decontracted(sent_1500)
        print(sent_1500)
        print("="*50)
```

was way to hot for my blood, took a bite and did a jig lol
=====

```
In [77]: #remove words with numbers python: https://stackoverflow.com/a/18082370/4084039
        sent_0 = re.sub("\S*\d\S*", "", sent_0).strip()
        print(sent_0)
```

My dogs loves this chicken but its a product from China, so we wont be buying it anymore. Its

```
In [78]: #remove spacial character: https://stackoverflow.com/a/5843547/4084039
        sent_1500 = re.sub('[^A-Za-z0-9]+', ' ', sent_1500)
        print(sent_1500)
```

was way to hot for my blood took a bite and did a jig lol

```
In [79]: # https://gist.github.com/sebleier/554280
        # we are removing the words from the stop words list: 'no', 'nor', 'not'
        # <br /><br /> ==> after the above steps, we are getting "br br"
        # we are including them into stop words list
        # instead of <br /> if we have <br/> these tags would have reuvmoved in the 1st step

        stopwords= set(['br', 'the', 'i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves',
                        "you'll", "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him',
                        'she', "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself',
                        'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', 'that',
                        'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had',
                        'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
                        'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into', 'through',
                        'above', 'below', 'to', 'from', 'up', 'down', 'in', 'out', 'on', 'off', 'over',
                        'then', 'once', 'here', 'there', 'when', 'where', 'why', 'how', 'all', 'any',
                        'most', 'other', 'some', 'such', 'only', 'own', 'same', 'so', 'than', 'too',
                        's', 't', 'can', 'will', 'just', 'don', "don't", 'should', "should've", 'n',
                        've', 'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't",
                        "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mi',
                        "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't",
                        'won', "won't", 'wouldn', "wouldn't"])
```

```
In [80]: # Combining all the above stundents
        from tqdm import tqdm
        preprocessed_reviews = []
        # tqdm is for printing the status bar
        for sentence in tqdm(final['Text'].values):
```

```

sentence = re.sub(r"http\S+", "", sentence)
sentence = BeautifulSoup(sentence, 'lxml').get_text()
sentence = decontracted(sentence)
sentence = re.sub("\S*\d\S*", "", sentence).strip()
sentence = re.sub('[^A-Za-z]+', ' ', sentence)
# https://gist.github.com/sebleier/554280
sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in stopwords)
preprocessed_reviews.append(sentence.strip())

```

100%|| 87773/87773 [00:32<00:00, 2708.34it/s]

In [81]: preprocessed_reviews[1500]

Out[81]: 'way hot blood took bite jig lol'

[3.2] Preprocessing Review Summary

In [82]: *## Similarly you can do preprocessing for review summary also.*

5 [4] Featurization

5.1 [4.1] BAG OF WORDS

```

In [83]: # #BoW
# count_vect = CountVectorizer() #in scikit-learn
# count_vect.fit(preprocessed_reviews)
# print("some feature names ", count_vect.get_feature_names()[:10])
# print('='*50)

# final_counts = count_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_counts))
# print("the shape of out text BOW vectorizer ",final_counts.get_shape())
# print("the number of unique words ", final_counts.get_shape()[1])

```

5.2 [4.2] Bi-Grams and n-Grams.

```

In [84]: # #bi-gram, tri-gram and n-gram

# #removing stop words like "not" should be avoided before building n-grams
# # count_vect = CountVectorizer(ngram_range=(1,2))
# # please do read the CountVectorizer documentation http://scikit-learn.org/stable/m

# # you can choose these numebrs min_df=10, max_features=5000, of your choice
# count_vect = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
# final_bigram_counts = count_vect.fit_transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_bigram_counts))
# print("the shape of out text BOW vectorizer ",final_bigram_counts.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_bigr

```

5.3 [4.3] TF-IDF

```
In [85]: # tf_idf_vect = TfidfVectorizer(ngram_range=(1,2), min_df=10)
# tf_idf_vect.fit(preprocessed_reviews)
# print("some sample features(unique words in the corpus)",tf_idf_vect.get_feature_names())
# print('='*50)

# final_tf_idf = tf_idf_vect.transform(preprocessed_reviews)
# print("the type of count vectorizer ",type(final_tf_idf))
# print("the shape of out text TFIDF vectorizer ",final_tf_idf.get_shape())
# print("the number of unique words including both unigrams and bigrams ", final_tf_idf.get_shape()[0])
```

5.4 [4.4] Word2Vec

```
In [86]: # # Using Google News Word2Vectors

# # in this project we are using a pretrained model by google
# # its 3.3G file, once you load this into your memory
# # it occupies ~9Gb, so please do this step only if you have >12G of ram
# # we will provide a pickle file wich contains a dict ,
# # and it contains all our courpus words as keys and model[word] as values
# # To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# # from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# # it's 1.9GB in size.

# # http://kavita-ganesan.com/gensim-word2vec-tutorial-starter-code/#.W17SRFAzZPY
# # you can comment this whole cell
# # or change these variable according to your need

# is_your_ram_gt_16g=False
# want_to_use_google_w2v = False
# want_to_train_w2v = True

# if want_to_train_w2v:
#     # min_count = 5 considers only words that occurred atleast 5 times
#     w2v_model=Word2Vec(list_of_sentence,min_count=5,size=50, workers=4)
#     print(w2v_model.wv.most_similar('great'))
#     print('='*50)
#     print(w2v_model.wv.most_similar('worst'))

# elif want_to_use_google_w2v and is_your_ram_gt_16g:
#     if os.path.isfile('GoogleNews-vectors-negative300.bin'):
#         w2v_model=KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin')
#         print(w2v_model.wv.most_similar('great'))
#         print(w2v_model.wv.most_similar('worst'))
#     else:
#         print("you don't have gogole's word2vec file, keep want_to_train_w2v = True")
```

```
In [87]: # w2v_words = list(w2v_model.wv.vocab)
# print("number of words that occurred minimum 5 times ", len(w2v_words))
# print("sample words ", w2v_words[0:50])
```

5.5 [4.4.1] Converting text into vectors using Avg W2V, TFIDF-W2V

[4.4.1.1] Avg W2v

```
In [88]: # # average Word2Vec
# # compute average word2vec for each review.
# sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
# for sent in tqdm(list_of_sentence): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length 50, you might need
#     cnt_words = 0; # num of words with a valid vector in the sentence/review
#     for word in sent: # for each word in a review/sentence
#         if word in w2v_words:
#             vec = w2v_model.wv[word]
#             sent_vec += vec
#             cnt_words += 1
#     if cnt_words != 0:
#         sent_vec /= cnt_words
#     sent_vectors.append(sent_vec)
# print(len(sent_vectors))
# print(len(sent_vectors[0]))
```

[4.4.1.2] TFIDF weighted W2v

```
In [89]: # # S = ["abc def pqr", "def def def abc", "pqr pqr def"]
# model = TfidfVectorizer()
# tf_idf_matrix = model.fit_transform(preprocessed_reviews)
# # we are converting a dictionary with word as a key, and the idf as a value
# dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))
```

```
In [90]: # # TF-IDF weighted Word2Vec
# tfidf_feat = model.get_feature_names() # tfidf words/col-names
# # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

# tfidf_sent_vectors = []; # the tfidf-w2v for each sentence/review is stored in this
# row=0;
# for sent in tqdm(list_of_sentence): # for each review/sentence
#     sent_vec = np.zeros(50) # as word vectors are of zero length
#     weight_sum = 0; # num of words with a valid vector in the sentence/review
#     for word in sent: # for each word in a review/sentence
#         if word in w2v_words and word in tfidf_feat:
#             vec = w2v_model.wv[word]
#             tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
#             # to reduce the computation we are
#             # dictionary[word] = idf value of word in whole corpus
#             # sent.count(word) = tf value of word in this review
```

```

#             tf_idf = dictionary[word]*(sent.count(word)/len(sent))
#             sent_vec += (vec * tf_idf)
#             weight_sum += tf_idf
#         if weight_sum != 0:
#             sent_vec /= weight_sum
#         tfidf_sent_vectors.append(sent_vec)
#         row += 1

```

6 [5] Assignment 9: Random Forests

Apply Random Forests & GBDT on these feature sets

- SET 1:** Review text, preprocessed one converted into vectors

- SET 2:** Review text, preprocessed one converted into vectors

- SET 3:** Review text, preprocessed one converted into vectors

- SET 4:** Review text, preprocessed one converted into vectors

The hyper paramter tuning (Consider two hyperparameters: n_estimators & max_depth)

- Find the best hyper parameter which will give the maximum <https://www.appliedaicom.com>

- Find the best hyper paramter using k-fold cross validation or simple cross validation data

- Use gridsearch cv or randomsearch cv or you can also write your own for loops to do this task

Feature importance

- Get top 20 important features and represent them in a word cloud. Do this for BOW & TFIDF.

Feature engineering

- To increase the performance of your model, you can also experiment with with feature engineering

-

- Taking length of reviews as another feature.

- Considering some features from review summary as well.

Representation of results

- You need to plot the performance of model both on train data and cross validation data for

 with X-axis as **n_estimators**, Y-axis as **score**

```

    <p style="text-align:center;font-size:30px;color:red;"><strong>(or)</strong></p> <br>
    <li>You need to plot the performance of model both on train data and cross validation data for
    <img src='heat_map.JPG' width=300px> <a href='https://seaborn.pydata.org/generated/seaborn.heat
    <li>You choose either of the plotting techniques out of 3d plot or heat map</li>
    <li>Once after you found the best hyper parameter, you need to train your model with it, and f
    <img src='train_test_auc.JPG' width=300px></li>
    <li>Along with plotting ROC curve, you need to print the <a href='https://www.appliedaicourse.
    <img src='confusion_matrix.png' width=300px></li>
    </ul>
</li>
<br>
<li><strong>Conclusion</strong>
    <ul>
<li>You need to summarize the results at the end of the notebook, summarize it in the table for
    <img src='summary.JPG' width=400px>
</li>
    </ul>

```

Note: Data Leakage

1. There will be an issue of data-leakage if you vectorize the entire data and then split it into train/cv/test.
2. To avoid the issue of data-leakag, make sure to split your data first and then vectorize it.
3. While vectorizing your data, apply the method fit_transform() on you train data, and apply the method transform() on cv/test data.
4. For more details please go through this link.

6.1 [5.1] Applying RF

```

In [91]: from sklearn.metrics import confusion_matrix
        from sklearn.metrics import roc_curve, auc
        from sklearn.metrics import roc_auc_score
        from sklearn.model_selection import train_test_split
        from sklearn.model_selection import StratifiedKFold
        from sklearn.ensemble import RandomForestClassifier
        import pprint
        import os.path
        import pickle
        import math

        import xgboost as xgb
        from wordcloud import WordCloud, STOPWORDS
        from PIL import Image
        import urllib
        import requests

        import warnings
        warnings.filterwarnings('ignore')

```

6.1.1 [5.0.0] Splitting up the Dataset into D_train and D_test

```
In [92]: num_data_points = 50000
```

```
In [93]: Dx_train, Dx_test, Dy_train, Dy_test = train_test_split(preprocessed_reviews[:num_data
```

```
In [94]: prettytable_data = []
```

6.1.2 [5.0.1] Defining some functions to increase code reusability and readability

```
In [95]: '''Creating Custom Vectorizers for TFIDF - W2Vec and Avg - W2Vec'''
```

```
class Tfidf_W2Vec_Vectorizer(object):
    def __init__(self, w2vec_model):
        if(w2v_model is None):
            raise Exception('Word 2 Vector model passed to Tfidf_W2Vec Vectorizer is None')
        self.tfidf = TfidfVectorizer(max_features=300)
        self.dictionary = None
        self.tfidf_feat = None

        self.word2vec = w2vec_model

    def fit(self, X):
        '''X : list'''
        #Initializing the TFIDF Vectorizer
        self.tfidf.fit_transform(X)
        # we are converting a dictionary with word as a key, and the idf as a value
        self.dictionary = dict(zip(self.tfidf.get_feature_names(), list(self.tfidf.idf)))
        self.tfidf_feat = self.tfidf.get_feature_names()

        return self

    def transform(self, X):
        '''X : list'''
        return np.array([
            np.mean([self.word2vec[w] * self.dictionary[word]*(X.count(word)/len(X))
                     for w in words if w in self.word2vec and w in self.tfidf_feat], axis=0)
            for words in X
        ])

class Avg_W2Vec_Vectorizer(object):
    def __init__(self, w2vec_model):
        if(w2v_model is None):
            raise Exception('Word 2 Vector model passed to Avg_W2Vec Vectorizer is None')
        self.word2vec = w2vec_model

    def fit(self, X):
        return self
```

```

def transform(self, X):
    '''X : list'''
    return np.array([
        np.mean([self.word2vec[w] for w in words if w in self.word2vec]
                or [np.zeros(300)], axis=0)
        for words in X
    ])

```

```

In [96]: def get_vectorizer(vectorizer, train, W2V_model=None):
    if(vectorizer=='BOW'):
        vectorizer = CountVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
    if(vectorizer=='TFIDF'):
        vectorizer = TfidfVectorizer(ngram_range=(1,2), min_df=10, max_features=5000)
    if(vectorizer=='TFIDF-W2Vec'):
        vectorizer = Tfidf_W2Vec_Vectorizer(W2V_model)
    if(vectorizer=='Avg-W2Vec'):
        vectorizer = Avg_W2Vec_Vectorizer(W2V_model)

    vectorizer.fit(train)
    return vectorizer

```

```

In [97]: '''Perform Simple Cross Validation'''
def perform_hyperparameter_tuning(X, Y, vectorizer, results_path, retrain=False, W2V_model=None):
    #If the pandas dataframe with the hyperparameter info exists then return it

    if(retrain==False):
        # If Cross Validation results exists then return them
        if(os.path.exists(results_path)):
            return pd.read_csv(results_path)
        else:
            # If no data exists but retrain=False then mention accordingly
            print('Retrain is set to be False but no Cross Validation Results DataFrame exists')
    else:
        # else perform hyperparameter tuning
        print('Performing Hyperparameter Tuning...\n')
        # regularization parameter
        hyperparameters = {
            'n_estimators' : [1, 20, 40, 60, 80, 100, 120, 140, 160, 180, 200]
        }

        n_estimators = []

        train_scores = []
        test_scores = []

        train_mean_score = []
        test_mean_score = []

```



```

# Initializing KFold
skf = StratifiedKFold(n_splits=3)
X = np.array(X)
Y = np.array(Y)

for estimators in hyperparameters['n_estimators']:

    #Performing Cross Validation
    for train_index, test_index in skf.split(X, Y):
        Dx_train, Dx_cv = X[train_index], X[test_index]
        Dy_train, Dy_cv = Y[train_index], Y[test_index]

        #Initializing the Vectorizer
        vectorizer = get_vectorizer(vectorizer, Dx_train.tolist(), W2V_model)

        #Transforming the data to features
        x_train = vectorizer.transform(Dx_train.tolist())
        x_cv = vectorizer.transform(Dx_cv.tolist())

        #Initializing the LR model
        clf = RandomForestClassifier(n_estimators=estimators, class_weight='b

        # Fit the model
        clf.fit(x_train, Dy_train)

        #Prediction
        train_results = clf.predict_proba(x_train)
        cv_results = clf.predict_proba(x_cv)

        try:
            train_score = roc_auc_score(Dy_train, train_results[:, 1])
            test_score = roc_auc_score(Dy_cv, cv_results[:, 1])

            #storing the results to form a dataframe
            train_scores.append(train_score)
            test_scores.append(test_score)

        except Exception as e:
            print('Error Case : ', e)
            print(('Actual, Predicted'))
            [print((Dy_cv[i], cv_results[i, 1])) for i in range(len(Dy_cv))]

        print('CV iteration : n_estimators={0}, train_score={1}, test_score={2}'.format(estimators, train_score, test_score))

    train_mean_score.append(sum(train_scores)/len(train_scores))
    test_mean_score.append(sum(test_scores)/len(test_scores))

```

```

n_estimators.append(estimators)

print('CV : n_estimators={0}, train_score={1}, test_score={2}'
      .format(estimators, sum(train_scores)/len(train_scores), sum(test_s

train_scores = []
test_scores = []

# Creating a DataFrame from the saved data for visualization
results_df = pd.DataFrame({
    'n_estimators' : n_estimators,
    'train_score' : train_mean_score,
    'test_score' : test_mean_score
})

#writing the results to csv after performing hyperparameter tuning
try:
    results_df.to_csv(results_path)
except Exception as ex:
    print(str(ex), "\nError occured while converting DataFrame to CSV after c
return results_df

```

```

In [98]: def analyse_results(df):
    # plotting error curves
    fig = plt.figure()
    ax = fig.gca()

    plt.plot([i for i in df.n_estimators.tolist()], df.test_score.tolist(), '-o', c='b')
    plt.plot([i for i in df.n_estimators.tolist()], df.train_score.tolist(), '-o', c='r')
    plt.grid(True)
    plt.xlabel('Hyperparameter : n_estimators')
    plt.ylabel('Area Under ROC Curve')
    plt.title('AUC ROC Curve for Random Forest Classifier')
    plt.legend(loc='best')
    plt.show()

    # return the best parameters
    mmax = 0
    ind_max = 0
    for index, row in df.iterrows():
        if(row['test_score']>mmax):
            mmax=row['test_score']
            ind_max = index

    best_params = {
        'n_estimators':df.loc[ind_max, 'n_estimators']
    }

```

```

        return best_params

In [99]: def retrain_with_best_params(data, labels, best_params, vec_name, model_path, word2vec):
    if os.path.exists(model_path):
        print('Loading Model....')
        with open(model_path, 'rb') as input_file:
            clf = pickle.load(input_file)
    else:
        clf = RandomForestClassifier(n_estimators=best_params['n_estimators'], class_weight='balanced')

        print('Initializing Vectorizer')
        vectorizer = get_vectorizer(vectorizer=vec_name, train=data, W2V_model=word2vec)
        print('Training Model....')
        clf.fit(vectorizer.transform(data), np.array(labels))

        print('Saving Trained Model....')
        with open(model_path, 'wb') as file:
            pickle.dump(clf, file)
    return clf

In [100]: def plot_confusion_matrix(model, data, labels, dataset_label):
    pred = model.predict(data)
    conf_mat = confusion_matrix(labels, pred)

    strings = strings = np.asarray([['TN = ', 'FP = '],
                                     ['FN = ', 'TP = ']])

    labels = (np.asarray(["{0}{1}".format(string, value)
                           for string, value in zip(strings.flatten(),
                                                       conf_mat.flatten())])
              ).reshape(2, 2)

    fig, ax = plt.subplots()
    ax.set(xlabel='Predicted', ylabel='Actual', title='Confusion Matrix : {0}'.format(dataset_label))
    sns.heatmap(conf_mat, annot=labels, fmt="", cmap='YlGnBu', ax=ax)
    ax.set_xlabel('Predicted')
    ax.set_ylabel('Actual')
    ax.set_xticklabels(['False', 'True'])
    ax.set_yticklabels(['False', 'True'])
    plt.show()

In [101]: def plot_AUC_ROC(model, vectorizer, Dx_train, Dx_test, Dy_train, Dy_test):
    #predicting probability of Dx_test, Dx_train
    test_score = model.predict_proba(vectorizer.transform(Dx_test))
    train_score = model.predict_proba(vectorizer.transform(Dx_train))

```

```

#Finding out the ROC_AUC_SCORE
train_roc_auc_score = roc_auc_score(np.array(Dy_train), train_score[:, 1])
print('Area Under the Curve for Train : ', train_roc_auc_score)
test_roc_auc_score = roc_auc_score(np.array(Dy_test), test_score[:, 1])
print('Area Under the Curve for Test : ', test_roc_auc_score)

#Plotting with matplotlib.pyplot
#ROC Curve for D-train
train_fpr, train_tpr, thresholds = roc_curve(np.array(Dy_train), train_score[:, 1])
plt.plot(train_fpr, train_tpr, label="train AUC =" + str(auc(train_fpr, train_tpr)))

# ROC Curve for D-test
test_fpr, test_tpr, thresholds = roc_curve(np.array(Dy_test), test_score[:, 1])
plt.plot(test_fpr, test_tpr, label="train AUC =" + str(auc(test_fpr, test_tpr)))

plt.legend()
plt.xlabel("FPR : False Positive Ratio")
plt.ylabel("TPF : True Positive Ratio")
plt.title("Area Under ROC Curve")
plt.show()

plot_confusion_matrix(model, vectorizer.transform(Dx_train), np.array(Dy_train),
plot_confusion_matrix(model, vectorizer.transform(Dx_test), np.array(Dy_test), '
return train_roc_auc_score, test_roc_auc_score

```

```

In [102]: def generate_wordcloud(words, mask):
    word_cloud = WordCloud(width = 512, height = 512, background_color='white', stopw
    plt.figure(figsize=(10,8),facecolor = 'white', edgecolor='blue')
    plt.imshow(word_cloud)
    plt.axis('off')
    plt.tight_layout(pad=0)
    plt.show()

```

6.1.3 [5.1.1] Applying Random Forests on BOW, SET 1

```

In [103]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/BOW_rf_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='BOW',
                                           results_path=csv_path, retrain=False, W2V

# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf.pkl'.format('BOW')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'BOW', model_path

print('Retraining Vectorizer with Dx_train')

```


```

vectorizer_obj = get_vectorizer(W2V_model = None, train=Dx_train, vectorizer='BOW')

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train)

# appending the data results
prettytable_data.append(['BOW', 'Random Forests', best_parameters['n_estimators'], train_score, test_score])



```




```

{'n_estimators': 180}
Loading Model...
Retraining Vectorizer with Dx_train
Area Under the Curve for Train : 0.9999967997907125
Area Under the Curve for Test : 0.927260543464937

```





6.1.4 [5.1.2] Wordcloud of top 20 important features from SET 1

6.1.5 Reference article

<https://blog.goodaudience.com/how-to-generate-a-word-cloud-of-any-shape-in-python-7bce27a55f6e>

```
In [104]: # Please write all the code with proper documentation
feature_names = vectorizer_obj.get_feature_names()
feature_importances = clf.feature_importances_
features_with_importances = [(feature_names[i], feature_importances[i]) for i in range(
    feature_names.shape[0])]

# sorting the features with importances
features_with_importances.sort(key = lambda x: abs(x[1]), reverse=True)
important_features = features_with_importances[:20]

mask = np.array(Image.open(requests.get('http://www.clker.com/cliparts/0/i/x/Y/q/P/y/0iXyqP/y0iXyqP.png').content).getdata())

words = ' '.join([important_features[i][0] for i in range(20)])
generate_wordcloud(words, mask)
```



6.1.6 [5.1.3] Applying Random Forests on TFIDF, SET 2

```
In [105]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/TFIDF_rf_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           results_path=csv_path, retrain=False, W2V=False)

# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf.pkl'.format('TFIDF')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF', model_path)

print('Retraining Vectorizer with Dx_train')
```




```

vectorizer_obj = get_vectorizer(W2V_model = None, train=Dx_train, vectorizer='TFIDF')

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train)

# appending the data results
prettytable_data.append(['TFIDF', 'Random Forests', best_parameters['n_estimators'],



```




```

{'n_estimators': 180}
Loading Model...
Retraining Vectorizer with Dx_train
Area Under the Curve for Train : 0.9999967997907125
Area Under the Curve for Test : 0.9294060485380355

```





6.1.7 [5.1.4] Wordcloud of top 20 important features from SET 2

```
In [106]: # Please write all the code with proper documentation
feature_names = vectorizer_obj.get_feature_names()
feature_importances = clf.feature_importances_
features_with_importances = [(feature_names[i], feature_importances[i]) for i in range(len(feature_names))]

# sorting the features with importances
features_with_importances.sort(key = lambda x: abs(x[1]), reverse=True)
important_features = features_with_importances[:20]

mask = np.array(Image.open(requests.get('http://www.clker.com/cliparts/O/i/x/Y/q/P/y').content))

words = ' '.join([important_features[i][0] for i in range(20)])
generate_wordcloud(words, mask)
```



6.2 Preparing/Training Google Word2Vec

```
In [107]: is_your_ram_gt_16g=True
          want_to_use_google_w2v = True
          want_to_train_w2v = False

          path_to_word2vec = '/home/monodeepdas112/Datasets/GoogleNews-vectors-negative300.bin

          if want_to_train_w2v:

              # Train your own Word2Vec model using your own text corpus
              i=0
              list_of_sentences=[]
              for sentence in preprocessed_reviews:
                  list_of_sentences.append(sentence.split())
```

```

# min_count = 5 considers only words that occurred at least 5 times
w2v_model=Word2Vec(list_of_sentences,min_count=5,size=300, workers=4)
print(w2v_model.wv.most_similar('great'))
print('='*50)
print(w2v_model.wv.most_similar('worst'))

elif want_to_use_google_w2v and is_your_ram_gt_16g:
    if os.path.isfile(path_to_word2vec):
        print('Preparing to load pre-trained Word2Vec model !')
        w2v_model=KeyedVectors.load_word2vec_format(path_to_word2vec, binary=True)
        print('Successfully loaded model into memory !!')
        print('Words similar to "similar" : ', w2v_model.wv.most_similar('great'))
        print('Words similar to "worst" : ',w2v_model.wv.most_similar('worst'))
    else:
        print("you don't have google's word2vec file, keep want_to_train_w2v = True,

```

Preparing to load pre-trained Word2Vec model !

Successfully loaded model into memory !!

Words similar to "similar" : [('terrific', 0.798933207988739), ('fantastic', 0.79352122545242)

Words similar to "worst" : [('Worst', 0.6146091222763062), ('weakest', 0.6143776774406433), (

6.2.1 [5.1.5] Applying Random Forests on AVG W2V, SET 3

In [108]: *# Please write all the code with proper documentation*

```
csv_path = 'saved_models/Assignment9/Avg-W2Vec_rf_results.csv'
```

```
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='Avg-W2Vec',
                                           results_path=csv_path, retrain=False, W2V_model=w2v_model)
```

```
# Analysing best parameters
```

```
best_parameters = analyse_results(cv_results)
```

```
pprint.pprint(best_parameters)
```

```
# retraining the model with best parameters
```

```
model_path = 'saved_models/Assignment9/{0}_rf.pkl'.format('Avg-W2Vec')
```

```
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'Avg-W2Vec', model_path)
```

```
print('Retraining Vectorizer with Dx_train')
```

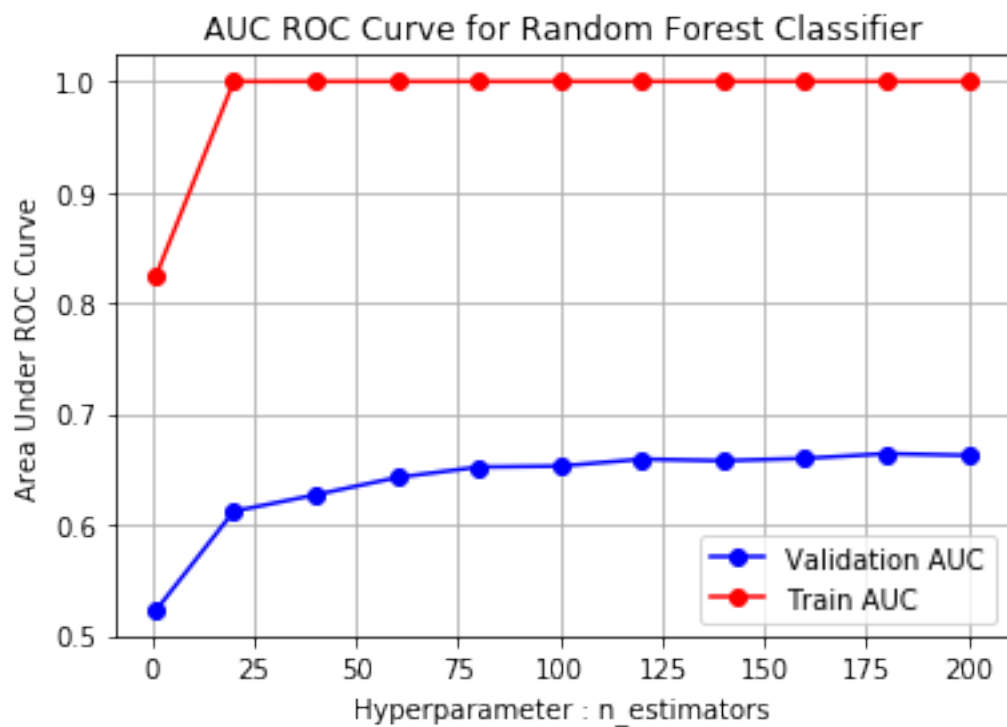
```
vectorizer_obj = get_vectorizer(W2V_model = w2v_model, train=Dx_train, vectorizer='Avg-W2Vec')
```

```
# plotting AUC ROC
```

```
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train)
```

```
# appending the data results
```

```
prettytable_data.append(['Avg-W2Vec', 'Random Forests', best_parameters['n_estimators'], train_score, test_score])
```





```
{'n_estimators': 180}
```


```
Loading Model...
```

```
Retraining Vectorizer with Dx_train
```

```
Area Under the Curve for Train : 0.9999969551406778
```

```
Area Under the Curve for Test : 0.6716041945109271
```





6.2.2 [5.1.6] Applying Random Forests on TFIDF W2V, SET 4

```
In [109]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/TFIDF-W2Vec-W2Vec_rf_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           results_path=csv_path, retrain=False, W2V_model=w2v_model)

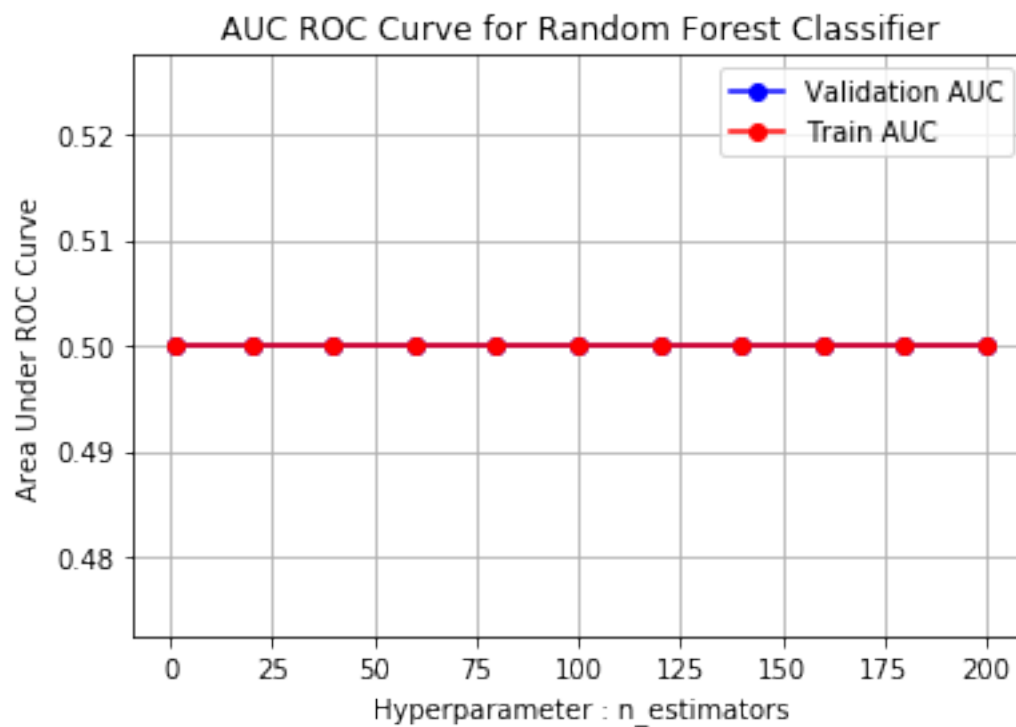
# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf.pkl'.format('TFIDF-W2Vec')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF-W2Vec', model_path=model_path)



print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = w2v_model, train=Dx_train, vectorizer='TFIDF')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train, Dy_test)

# appending the data results
prettytable_data.append(['TFIDF-W2Vec', 'Random Forests', best_parameters['n_estimators']])
```

```
{'n_estimators': 1}  
Loading Model...  
Retraining Vectorizer with Dx_train  
Area Under the Curve for Train : 0.5  
Area Under the Curve for Test : 0.5
```





6.3 [5.2] Applying GBDT using XGBOOST

```
In [110]: '''Perform Simple Cross Validation'''
def perform_hyperparameter_tuning(X, Y, vectorizer, results_path, retrain=False, W2V):
    #If the pandas dataframe with the hyperparameter info exists then return it

    if(retrain==False):
        # If Cross Validation results exists then return them
        if(os.path.exists(results_path)):
            return pd.read_csv(results_path)
        else:
            # If no data exists but retrain=False then mention accordingly
            print('Retrain is set to be False but no Cross Validation Results DataFr

    else:
        # else perform hyperparameter tuning
        print('Performing Hyperparameter Tuning...\n')
        # regularization parameter
        hyperparameters = {
            'n_estimators' : [1, 20, 40, 60, 80, 100, 120, 140, 160, 180, 200]
        }

        n_estimators = []
```

```

train_scores = []
test_scores = []

train_mean_score = []
test_mean_score = []

# Initializing KFold
skf = StratifiedKFold(n_splits=3)
X = np.array(X)
Y = np.array(Y)

for estimators in hyperparameters['n_estimators']:

    #Performing Cross Validation
    for train_index, test_index in skf.split(X, Y):
        Dx_train, Dx_cv = X[train_index], X[test_index]
        Dy_train, Dy_cv = Y[train_index], Y[test_index]

        #Initializing the Vectorizer
        vectorizer = get_vectorizer(vectorizer, Dx_train.tolist(), W2V_model)

        #Transforming the data to features
        x_train = vectorizer.transform(Dx_train.tolist())
        x_cv = vectorizer.transform(Dx_cv.tolist())

        #Initializing the LR model
        clf = xgb.XGBClassifier(max_depth=5, n_estimators=estimators, colsamp

        # Fit the model
        clf.fit(x_train, Dy_train)

        #Prediction
        train_results = clf.predict_proba(x_train)
        cv_results = clf.predict_proba(x_cv)

    try:
        train_score = roc_auc_score(Dy_train, train_results[:, 1])
        test_score = roc_auc_score(Dy_cv, cv_results[:, 1])

        #storing the results to form a dataframe
        train_scores.append(train_score)
        test_scores.append(test_score)

    except Exception as e:
        print('Error Case : ', e)
        print(('Actual, Predicted'))
        [print((Dy_cv[i], cv_results[i, 1])) for i in range(len(Dy_cv))]

```

```

        print('CV iteration : n_estimators={0}, train_score={1}, test_score={2}'
              .format(estimators, train_score, test_score))

    train_mean_score.append(sum(train_scores)/len(train_scores))
    test_mean_score.append(sum(test_scores)/len(test_scores))

    n_estimators.append(estimators)

    print('CV : n_estimators={0}, train_score={1}, test_score={2}'
          .format(estimators, sum(train_scores)/len(train_scores), sum(test_scores)/len(test_scores)))

    train_scores = []
    test_scores = []

    # Creating a DataFrame from the saved data for visualization
    results_df = pd.DataFrame({
        'n_estimators' : n_estimators,
        'train_score' : train_mean_score,
        'test_score' : test_mean_score
    })

    #writing the results to csv after performing hyperparameter tuning
    try:
        results_df.to_csv(results_path)
    except Exception as ex:
        print(str(ex), "\nError occured while converting DataFrame to CSV after hyperparameter tuning")
    return results_df

```

6.3.1 [5.2.1] Applying XGBOOST on BOW, SET 1

```

In [111]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/BOW_rf_xgb_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='BOW',
                                           results_path=csv_path, retrain=False, W2V=W2V_model)

# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)


# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf_xgb.pkl'.format('BOW')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'BOW', model_path)

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = None, train=Dx_train, vectorizer='BOW')



# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train, Dy_test)


```

```
# appending the data results
prettytable_data.append(['BOW', 'Random Forests XGB', best_parameters['n_estimators']])
```



```
{'n_estimators': 200}
Loading Model...
Retraining Vectorizer with Dx_train
Area Under the Curve for Train : 0.9999967997907124
Area Under the Curve for Test : 0.9273981165911194
```





6.3.2 [5.2.2] Applying XGBOOST on TFIDF, SET 2

```
In [112]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/TFIDF_rf_xgb_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           results_path=csv_path, retrain=False, W2V_model=W2V_model)

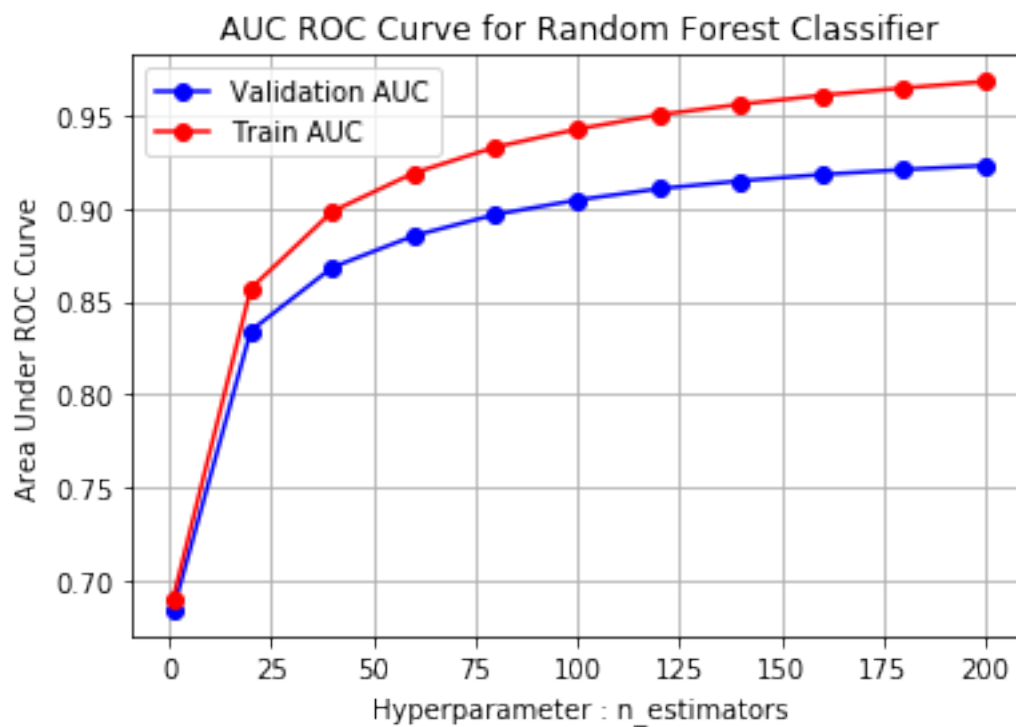
# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf_xgb.pkl'.format('TFIDF')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF', model_path)

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = None, train=Dx_train, vectorizer='TFIDF')

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train, Dy_test)

# appending the data results
prettytable_data.append(['TFIDF', 'Random Forests XGB', best_parameters['n_estimators']])
```



```
{'n_estimators': 200}
```


```
Loading Model...
```

```
Retraining Vectorizer with Dx_train
```

```
Area Under the Curve for Train : 0.9999967997907124
```

```
Area Under the Curve for Test : 0.9308974530549805
```





6.3.3 [5.2.3] Applying XGBOOST on AVG W2V, SET 3

```
In [113]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/Avg-W2Vec_rf_xgb_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='Avg-W2Vec',
                                           results_path=csv_path, retrain=False, W2V_model=w2v_model)


# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf_xgb.pkl'.format('Avg-W2Vec')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'Avg-W2Vec', model_path=model_path)

print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = w2v_model, train=Dx_train, vectorizer='Avg-W2Vec')

# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train, Dy_test)

# appending the data results
prettytable_data.append(['Avg-W2Vec', 'Random Forests XGB', best_parameters['n_estimators']])
```





```
{'n_estimators': 100}
```


```
Loading Model...
```

```
Retraining Vectorizer with Dx_train
```

```
Area Under the Curve for Train : 0.9999969551406779
```

```
Area Under the Curve for Test : 0.6652490383110953
```





6.3.4 [5.2.4] Applying XGBOOST on TFIDF W2V, SET 4

```
In [114]: # Please write all the code with proper documentation
csv_path = 'saved_models/Assignment9/TFIDF-W2Vec_rf_xgb_results.csv'
cv_results = perform_hyperparameter_tuning(X=Dx_train, Y=Dy_train, vectorizer='TFIDF',
                                           results_path=csv_path, retrain=False, W2V_model=w2v_model)


# Analysing best parameters
best_parameters = analyse_results(cv_results)
pprint.pprint(best_parameters)

# retraining the model with best parameters
model_path = 'saved_models/Assignment9/{0}_rf_xgb.pkl'.format('TFIDF-W2Vec')
clf = retrain_with_best_params(Dx_train, Dy_train, best_parameters, 'TFIDF-W2Vec', model_path=model_path)



print('Retraining Vectorizer with Dx_train')
vectorizer_obj = get_vectorizer(W2V_model = w2v_model, train=Dx_train, vectorizer='TFIDF')


# plotting AUC ROC
train_score, test_score = plot_AUC_ROC(clf, vectorizer_obj, Dx_train, Dx_test, Dy_train, Dy_test)

# appending the data results
prettytable_data.append(['TFIDF-W2Vec', 'Random Forests XGB', best_parameters['n_estimators']])
```



```
{'n_estimators': 1}  
Loading Model...  
Retraining Vectorizer with Dx_train  
Area Under the Curve for Train : 0.5  
Area Under the Curve for Test : 0.5
```





7 [6] Conclusions

```
In [119]: from prettytable import PrettyTable
           # Please compare all your models using Prettytable library
           x = PrettyTable()

           x.field_names = ["Vectorizer", "Model", "n_estimators", "Train AUC", "Test AUC"]
           [x.add_row(i) for i in prettytable_data]
           print(x)
```

Vectorizer	Model	n_estimators	Train AUC	Test AUC
BOW	Random Forests	180	0.9999967997907125	0.927260543464937
TFIDF	Random Forests	180	0.9999967997907125	0.9294060485380355
Avg-W2Vec	Random Forests	180	0.9999969551406778	0.6716041945109271
TFIDF-W2Vec	Random Forests	1	0.5	0.5
BOW	Random Forests XGB	200	0.9999967997907124	0.9273981165911194
TFIDF	Random Forests XGB	200	0.9999967997907124	0.9308974530549805
Avg-W2Vec	Random Forests XGB	100	0.9999969551406779	0.6652490383110953
TFIDF-W2Vec	Random Forests XGB	1	0.5	0.5